

3_RT_Models

June 5, 2024

1 Machine Learning Models Applied to the IDS-2017

The purpose of this notebook is to experiment different machine learning on the IDS-2017 dataset generated by the CICFlowMeter on the recorder traffic which included benign and malicious flows

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import glob
import os
import xgboost as xgb
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, average_precision_score, \
    make_scorer, precision_score, accuracy_score, confusion_matrix
from notebook_utils import upsample_dataset
%matplotlib inline
%load_ext autoreload
%autoreload 2
file_path = \
    r"CIC-IDS-2017\CSVs\GeneratedLabelledFlows\TrafficLabelling\processed\ids2017_processed.\
    CSV"

def replace_invalid(df):
    # Select only numeric columns
    numeric_columns = df.select_dtypes(include=[np.number]).columns

    # Identify columns with NaN, infinity, or negative values
    invalid_columns = df[numeric_columns].columns[df[numeric_columns].isna().
        any() |
        np.isinf(df[numeric_columns])).
        any() |
        (df[numeric_columns] < 0).
        any()]
```

```

    print("Columns with NaN, infinity, or negative values:", invalid_columns.
    ↪tolist())

    # Replace invalid values with NaN and fill with column mean
    df[invalid_columns] = df[invalid_columns].replace([np.inf, -np.inf, -1], np.
    ↪nan)
    df[invalid_columns] = df[invalid_columns].fillna(df[invalid_columns].mean())

    return df

def load_dataset(file_path):
    df = pd.read_csv(file_path)
    convert_dict = {'label': 'category'}
    df = df.astype(convert_dict)
    replace_invalid(df)
    df.info()
    return df

attack_labels = {
    0: 'BENIGN',
    7: 'FTP-Patator',
    11: 'SSH-Patator',
    6: 'DoS slowloris',
    5: 'DoS Slowhttptest',
    4: 'DoS Hulk',
    3: 'DoS GoldenEye',
    8: 'Heartbleed',
    12: 'Web Attack - Brute Force',
    14: 'Web Attack - XSS',
    13: 'Web Attack - Sql Injection',
    9: 'Infiltration',
    1: 'Bot',
    10: 'PortScan',
    2: 'DDoS'
}

```

2 1. Preparing the Dataset

```
[2]: df = load_dataset(file_path)
```

```

Columns with NaN, infinity, or negative values: ['flow_duration',
'flow_bytes_s', 'flow_packets_s', 'flow_iat_mean', 'flow_iat_max',
'flow_iat_min', 'fwd_iat_min', 'fwd_header_length', 'bwd_header_length',
'fwd_header_length_1', 'init_win_bytes_forward', 'init_win_bytes_backward',
'min_seg_size_forward']
<class 'pandas.core.frame.DataFrame'>

```

RangeIndex: 2830743 entries, 0 to 2830742

Data columns (total 96 columns):

#	Column	Dtype
0	destination_port	int64
1	protocol	int64
2	flow_duration	float64
3	total_fwd_packets	int64
4	total_backward_packets	int64
5	total_length_of_fwd_packets	float64
6	total_length_of_bwd_packets	float64
7	fwd_packet_length_max	float64
8	fwd_packet_length_min	float64
9	fwd_packet_length_mean	float64
10	fwd_packet_length_std	float64
11	bwd_packet_length_max	float64
12	bwd_packet_length_min	float64
13	bwd_packet_length_mean	float64
14	bwd_packet_length_std	float64
15	flow_bytes_s	float64
16	flow_packets_s	float64
17	flow_iat_mean	float64
18	flow_iat_std	float64
19	flow_iat_max	float64
20	flow_iat_min	float64
21	fwd_iat_total	float64
22	fwd_iat_mean	float64
23	fwd_iat_std	float64
24	fwd_iat_max	float64
25	fwd_iat_min	float64
26	bwd_iat_total	float64
27	bwd_iat_mean	float64
28	bwd_iat_std	float64
29	bwd_iat_max	float64
30	bwd_iat_min	float64
31	fwd_psh_flags	int64
32	bwd_psh_flags	int64
33	fwd_urg_flags	int64
34	bwd_urg_flags	int64
35	fwd_header_length	int64
36	bwd_header_length	int64
37	fwd_packets_s	float64
38	bwd_packets_s	float64
39	min_packet_length	float64
40	max_packet_length	float64
41	packet_length_mean	float64
42	packet_length_std	float64
43	packet_length_variance	float64

44	fin_flag_count	int64
45	syn_flag_count	int64
46	rst_flag_count	int64
47	psh_flag_count	int64
48	ack_flag_count	int64
49	urg_flag_count	int64
50	cwe_flag_count	int64
51	ece_flag_count	int64
52	down_up_ratio	float64
53	average_packet_size	float64
54	avg_fwd_segment_size	float64
55	avg_bwd_segment_size	float64
56	fwd_header_length_1	int64
57	fwd_avg_bytes_bulk	int64
58	fwd_avg_packets_bulk	int64
59	fwd_avg_bulk_rate	int64
60	bwd_avg_bytes_bulk	int64
61	bwd_avg_packets_bulk	int64
62	bwd_avg_bulk_rate	int64
63	subflow_fwd_packets	int64
64	subflow_fwd_bytes	int64
65	subflow_bwd_packets	int64
66	subflow_bwd_bytes	int64
67	init_win_bytes_forward	float64
68	init_win_bytes_backward	float64
69	act_data_pkt_fwd	int64
70	min_seg_size_forward	float64
71	active_mean	float64
72	active_std	float64
73	active_max	float64
74	active_min	float64
75	idle_mean	float64
76	idle_std	float64
77	idle_max	float64
78	idle_min	float64
79	label	category
80	is_attack	int64
81	label_code	int64
82	is_dos_hulk	int64
83	is_portscan	int64
84	is_ddos	int64
85	is_dos_goldeneye	int64
86	is_ftppatator	int64
87	is_sshpatator	int64
88	is_dos_slowloris	int64
89	is_dos_slowhttptest	int64
90	is_bot	int64
91	is_web_attack_brute_force	int64

```

92 is_web_attack_xss          int64
93 is_infiltration            int64
94 is_web_attack_sql_injection int64
95 is_heartbleed              int64
dtypes: category(1), float64(49), int64(46)
memory usage: 2.0 GB

```

```

[3]: X = df.iloc[:, 0:79]
      Y = df.iloc[:, 79:]
      X.info()
      Y.info()
      print(Y.label.value_counts())

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2830743 entries, 0 to 2830742
Data columns (total 79 columns):

```

#	Column	Dtype
0	destination_port	int64
1	protocol	int64
2	flow_duration	float64
3	total_fwd_packets	int64
4	total_backward_packets	int64
5	total_length_of_fwd_packets	float64
6	total_length_of_bwd_packets	float64
7	fwd_packet_length_max	float64
8	fwd_packet_length_min	float64
9	fwd_packet_length_mean	float64
10	fwd_packet_length_std	float64
11	bwd_packet_length_max	float64
12	bwd_packet_length_min	float64
13	bwd_packet_length_mean	float64
14	bwd_packet_length_std	float64
15	flow_bytes_s	float64
16	flow_packets_s	float64
17	flow_iat_mean	float64
18	flow_iat_std	float64
19	flow_iat_max	float64
20	flow_iat_min	float64
21	fwd_iat_total	float64
22	fwd_iat_mean	float64
23	fwd_iat_std	float64
24	fwd_iat_max	float64
25	fwd_iat_min	float64
26	bwd_iat_total	float64
27	bwd_iat_mean	float64
28	bwd_iat_std	float64
29	bwd_iat_max	float64

30	bwd_iat_min	float64
31	fwd_psh_flags	int64
32	bwd_psh_flags	int64
33	fwd_urg_flags	int64
34	bwd_urg_flags	int64
35	fwd_header_length	int64
36	bwd_header_length	int64
37	fwd_packets_s	float64
38	bwd_packets_s	float64
39	min_packet_length	float64
40	max_packet_length	float64
41	packet_length_mean	float64
42	packet_length_std	float64
43	packet_length_variance	float64
44	fin_flag_count	int64
45	syn_flag_count	int64
46	rst_flag_count	int64
47	psh_flag_count	int64
48	ack_flag_count	int64
49	urg_flag_count	int64
50	cwe_flag_count	int64
51	ece_flag_count	int64
52	down_up_ratio	float64
53	average_packet_size	float64
54	avg_fwd_segment_size	float64
55	avg_bwd_segment_size	float64
56	fwd_header_length_1	int64
57	fwd_avg_bytes_bulk	int64
58	fwd_avg_packets_bulk	int64
59	fwd_avg_bulk_rate	int64
60	bwd_avg_bytes_bulk	int64
61	bwd_avg_packets_bulk	int64
62	bwd_avg_bulk_rate	int64
63	subflow_fwd_packets	int64
64	subflow_fwd_bytes	int64
65	subflow_bwd_packets	int64
66	subflow_bwd_bytes	int64
67	init_win_bytes_forward	float64
68	init_win_bytes_backward	float64
69	act_data_pkt_fwd	int64
70	min_seg_size_forward	float64
71	active_mean	float64
72	active_std	float64
73	active_max	float64
74	active_min	float64
75	idle_mean	float64
76	idle_std	float64
77	idle_max	float64

```

78  idle_min                                float64
dtypes: float64(49), int64(30)
memory usage: 1.7 GB
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2830743 entries, 0 to 2830742
Data columns (total 17 columns):
 #   Column                                Dtype
---  -
 0   label                                category
 1   is_attack                           int64
 2   label_code                          int64
 3   is_dos_hulk                         int64
 4   is_portscan                        int64
 5   is_ddos                            int64
 6   is_dos_goldeneye                   int64
 7   is_ftppatator                      int64
 8   is_sshpatator                      int64
 9   is_dos_slowloris                   int64
10  is_dos_slowhttpptest                int64
11  is_bot                              int64
12  is_web_attack_brute_force           int64
13  is_web_attack_xss                   int64
14  is_infiltration                     int64
15  is_web_attack_sql_injection         int64
16  is_heartbleed                       int64
dtypes: category(1), int64(16)
memory usage: 348.3 MB
label
BENIGN                                2273097
DoS Hulk                             231073
PortScan                             158930
DDoS                                  128027
DoS GoldenEye                        10293
FTP-Patator                          7938
SSH-Patator                          5897
DoS slowloris                        5796
DoS Slowhttpptest                    5499
Bot                                   1966
Web Attack - Brute Force              1507
Web Attack - XSS                      652
Infiltration                          36
Web Attack - Sql Injection            21
Heartbleed                           11
Name: count, dtype: int64

```

3 2. Feature Selection

3.0.1 Correlation based feature selection

First, the columns with no variance are dropped as they have no impact on the target variables.

```
[4]: stats = X.describe()
std = stats.loc["std"]
features_no_var = std[std == 0.0].index
# Exclude non-numeric columns (e.g., categorical columns) from the features_
↳with zero variance
features_no_var_numeric = [col for col in features_no_var if col in X.
↳select_dtypes(include=[np.number]).columns]
print(features_no_var_numeric)
```

```
['bwd_psh_flags', 'bwd_urg_flags', 'fwd_avg_bytes_bulk', 'fwd_avg_packets_bulk',
'fwd_avg_bulk_rate', 'bwd_avg_bytes_bulk', 'bwd_avg_packets_bulk',
'bwd_avg_bulk_rate']
```

The destination port feature is dropped because it can act as a shortcut predictor and cause high overfitting for the training set as show in this [paper](#)

```
[5]: X = X.drop(columns=features_no_var)
X = X.drop(columns=['destination_port'])
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2830743 entries, 0 to 2830742
Data columns (total 70 columns):
#   Column                                Dtype
---  -
0   protocol                             int64
1   flow_duration                         float64
2   total_fwd_packets                     int64
3   total_backward_packets                int64
4   total_length_of_fwd_packets           float64
5   total_length_of_bwd_packets           float64
6   fwd_packet_length_max                  float64
7   fwd_packet_length_min                  float64
8   fwd_packet_length_mean                 float64
9   fwd_packet_length_std                  float64
10  bwd_packet_length_max                   float64
11  bwd_packet_length_min                   float64
12  bwd_packet_length_mean                  float64
13  bwd_packet_length_std                   float64
14  flow_bytes_s                           float64
15  flow_packets_s                         float64
16  flow_iat_mean                          float64
17  flow_iat_std                           float64
18  flow_iat_max                           float64
```


19	flow_iat_min	float64
20	fwd_iat_total	float64
21	fwd_iat_mean	float64
22	fwd_iat_std	float64
23	fwd_iat_max	float64
24	fwd_iat_min	float64
25	bwd_iat_total	float64
26	bwd_iat_mean	float64
27	bwd_iat_std	float64
28	bwd_iat_max	float64
29	bwd_iat_min	float64
30	fwd_psh_flags	int64
31	fwd_urg_flags	int64
32	fwd_header_length	int64
33	bwd_header_length	int64
34	fwd_packets_s	float64
35	bwd_packets_s	float64
36	min_packet_length	float64
37	max_packet_length	float64
38	packet_length_mean	float64
39	packet_length_std	float64
40	packet_length_variance	float64
41	fin_flag_count	int64
42	syn_flag_count	int64
43	rst_flag_count	int64
44	psh_flag_count	int64
45	ack_flag_count	int64
46	urg_flag_count	int64
47	cwe_flag_count	int64
48	ece_flag_count	int64
49	down_up_ratio	float64
50	average_packet_size	float64
51	avg_fwd_segment_size	float64
52	avg_bwd_segment_size	float64
53	fwd_header_length_1	int64
54	subflow_fwd_packets	int64
55	subflow_fwd_bytes	int64
56	subflow_bwd_packets	int64
57	subflow_bwd_bytes	int64
58	init_win_bytes_forward	float64
59	init_win_bytes_backward	float64
60	act_data_pkt_fwd	int64
61	min_seg_size_forward	float64
62	active_mean	float64
63	active_std	float64
64	active_max	float64
65	active_min	float64
66	idle_mean	float64

```

67 idle_std float64
68 idle_max float64
69 idle_min float64
dtypes: float64(49), int64(21)
memory usage: 1.5 GB

```

3.0.2 Remove collinear variables

```

[6]: threshold = 0.9
     corr_matrix = X.corr().abs()
     corr_matrix.head()

```

```

[6]:
protocol  flow_duration  total_fwd_packets \
protocol      1.000000      0.265288      0.007272
flow_duration  0.265288      1.000000      0.020857
total_fwd_packets  0.007272      0.020857      1.000000
total_backward_packets  0.006361      0.019669      0.999070
total_length_of_fwd_packets  0.033234      0.065456      0.365508

                                total_backward_packets \
protocol                                0.006361
flow_duration                            0.019669
total_fwd_packets                        0.999070
total_backward_packets                    1.000000
total_length_of_fwd_packets                0.359451

                                total_length_of_fwd_packets \
protocol                                0.033234
flow_duration                            0.065456
total_fwd_packets                        0.365508
total_backward_packets                    0.359451
total_length_of_fwd_packets                1.000000

                                total_length_of_bwd_packets \
protocol                                0.005191
flow_duration                            0.016186
total_fwd_packets                        0.996993
total_backward_packets                    0.994429
total_length_of_fwd_packets                0.353762

                                fwd_packet_length_max  fwd_packet_length_min \
protocol                                0.166066      0.315250
flow_duration                            0.273304      0.105235
total_fwd_packets                        0.009358      0.002989
total_backward_packets                    0.009039      0.002600
total_length_of_fwd_packets                0.197030      0.000275

                                fwd_packet_length_mean  fwd_packet_length_std \

```

protocol	0.052344	0.178832
flow_duration	0.143685	0.234434
total_fwd_packets	0.000032	0.001403
total_backward_packets	0.000333	0.001026
total_length_of_fwd_packets	0.185262	0.159787

	...	act_data_pkt_fwd	min_seg_size_forward	\
protocol	...	0.005043	0.003451	
flow_duration	...	0.015942	0.001357	
total_fwd_packets	...	0.887387	0.000184	
total_backward_packets	...	0.882566	0.000018	
total_length_of_fwd_packets	...	0.407448	0.001209	

	active_mean	active_std	active_max	active_min	\
protocol	0.085598	0.081018	0.109356	0.063663	
flow_duration	0.189298	0.241059	0.294033	0.121169	
total_fwd_packets	0.039937	0.008329	0.030459	0.041283	
total_backward_packets	0.038963	0.006437	0.028602	0.041278	
total_length_of_fwd_packets	0.101084	0.103326	0.126493	0.068325	

	idle_mean	idle_std	idle_max	idle_min
protocol	0.179676	0.071305	0.184514	0.170531
flow_duration	0.768031	0.243153	0.779524	0.738325
total_fwd_packets	0.001820	0.000809	0.001906	0.001670
total_backward_packets	0.001425	0.000492	0.001456	0.001330
total_length_of_fwd_packets	0.022660	0.027064	0.026079	0.018634

[5 rows x 70 columns]

```
[7]: # Upper triangle of correlations
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool))
upper.head()
```

```
[7]:
```

	protocol	flow_duration	total_fwd_packets	\
protocol	NaN	0.265288	0.007272	
flow_duration	NaN	NaN	0.020857	
total_fwd_packets	NaN	NaN	NaN	
total_backward_packets	NaN	NaN	NaN	
total_length_of_fwd_packets	NaN	NaN	NaN	

	total_backward_packets	\
protocol	0.006361	
flow_duration	0.019669	
total_fwd_packets	0.999070	
total_backward_packets	NaN	
total_length_of_fwd_packets	NaN	

	total_length_of_fwd_packets \
protocol	0.033234
flow_duration	0.065456
total_fwd_packets	0.365508
total_backward_packets	0.359451
total_length_of_fwd_packets	NaN

	total_length_of_bwd_packets \
protocol	0.005191
flow_duration	0.016186
total_fwd_packets	0.996993
total_backward_packets	0.994429
total_length_of_fwd_packets	0.353762

	fwd_packet_length_max	fwd_packet_length_min \
protocol	0.166066	0.315250
flow_duration	0.273304	0.105235
total_fwd_packets	0.009358	0.002989
total_backward_packets	0.009039	0.002600
total_length_of_fwd_packets	0.197030	0.000275

	fwd_packet_length_mean	fwd_packet_length_std \
protocol	0.052344	0.178832
flow_duration	0.143685	0.234434
total_fwd_packets	0.000032	0.001403
total_backward_packets	0.000333	0.001026
total_length_of_fwd_packets	0.185262	0.159787

	... act_data_pkt_fwd	min_seg_size_forward \
protocol	... 0.005043	0.003451
flow_duration	... 0.015942	0.001357
total_fwd_packets	... 0.887387	0.000184
total_backward_packets	... 0.882566	0.000018
total_length_of_fwd_packets	... 0.407448	0.001209

	active_mean	active_std	active_max	active_min \
protocol	0.085598	0.081018	0.109356	0.063663
flow_duration	0.189298	0.241059	0.294033	0.121169
total_fwd_packets	0.039937	0.008329	0.030459	0.041283
total_backward_packets	0.038963	0.006437	0.028602	0.041278
total_length_of_fwd_packets	0.101084	0.103326	0.126493	0.068325

	idle_mean	idle_std	idle_max	idle_min
protocol	0.179676	0.071305	0.184514	0.170531
flow_duration	0.768031	0.243153	0.779524	0.738325
total_fwd_packets	0.001820	0.000809	0.001906	0.001670
total_backward_packets	0.001425	0.000492	0.001456	0.001330

```
total_length_of_fwd_packets    0.022660    0.027064    0.026079    0.018634
```

```
[5 rows x 70 columns]
```

```
[8]: to_drop = [column for column in upper.columns if any(upper[column] > threshold)]
to_keep = [
    'Destination Port', 'Fwd Packet Length Std', 'Min Packet Length',
    'Packet Length Variance', 'PSH Flag Count', 'Active Max'
]
to_drop = [column for column in to_drop if column not in to_keep]
print('There are %d columns to remove.' % (len(to_drop)))
X = X.drop(columns=to_drop)
X.info()
```

There are 31 columns to remove.

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 2830743 entries, 0 to 2830742

Data columns (total 39 columns):

#	Column	Dtype
0	protocol	int64
1	flow_duration	float64
2	total_fwd_packets	int64
3	total_length_of_fwd_packets	float64
4	fwd_packet_length_max	float64
5	fwd_packet_length_min	float64
6	fwd_packet_length_mean	float64
7	bwd_packet_length_max	float64
8	bwd_packet_length_min	float64
9	flow_bytes_s	float64
10	flow_packets_s	float64
11	flow_iat_mean	float64
12	flow_iat_std	float64
13	flow_iat_min	float64
14	fwd_iat_min	float64
15	bwd_iat_total	float64
16	bwd_iat_mean	float64
17	bwd_iat_std	float64
18	bwd_iat_max	float64
19	fwd_psh_flags	int64
20	fwd_urg_flags	int64
21	fwd_header_length	int64
22	bwd_header_length	int64
23	bwd_packets_s	float64
24	min_packet_length	float64
25	fin_flag_count	int64
26	rst_flag_count	int64
27	psh_flag_count	int64

```

28  ack_flag_count          int64
29  urg_flag_count          int64
30  down_up_ratio           float64
31  init_win_bytes_forward  float64
32  init_win_bytes_backward float64
33  act_data_pkt_fwd        int64
34  min_seg_size_forward    float64
35  active_mean              float64
36  active_std               float64
37  active_max              float64
38  idle_std                 float64
dtypes: float64(27), int64(12)
memory usage: 842.3 MB

```

3.0.3 3. Split Dataset

The dataset is split into a train, cross-validation and evaluation sets with a ratio of 0.8/0.1/0.1. The dataset is stratified according to the label to have an equal representation of all classes in the 3 subsets.

```

[9]: X_train, X_temp, Y_train, Y_temp = train_test_split(X, Y, test_size=0.2,
↳stratify=Y.label_code)
X_eval, X_test, Y_eval, Y_test = train_test_split(X_temp, Y_temp, test_size=0.
↳5, stratify=Y_temp.label_code)

```

Upsampling is used to generate artificial samples for types of attacks that are underrepresented in the dataset.

```

[10]: X_train, Y_train = upsample_dataset(X, Y, 100000, attack_labels)

```

```

[11]: Y_train.label.value_counts()

```

```

[11]: label
BENIGN                2273097
DoS Hulk               231073
PortScan              158930
DDoS                  128027
Bot                   100000
DoS GoldenEye         100000
DoS Slowhttptest      100000
DoS slowloris         100000
FTP-Patator           100000
Heartbleed            100000
Infiltration          100000
SSH-Patator           100000
Web Attack - Brute Force 100000
Web Attack - Sql Injection 100000
Web Attack - XSS       100000

```

Name: count, dtype: int64

```
[12]: Y_eval.label.value_counts()
```

```
[12]: label
      BENIGN                227310
      DoS Hulk              23107
      PortScan             15893
      DDoS                 12803
      DoS GoldenEye        1029
      FTP-Patator           794
      SSH-Patator           589
      DoS slowloris        579
      DoS Slowhttptest     550
      Bot                  197
      Web Attack - Brute Force 151
      Web Attack - XSS       65
      Infiltration           4
      Web Attack - Sql Injection 2
      Heartbleed             1
      Name: count, dtype: int64
```

```
[13]: Y_test.label.value_counts()
```

```
[13]: label
      BENIGN                227310
      DoS Hulk              23108
      PortScan             15893
      DDoS                 12803
      DoS GoldenEye        1030
      FTP-Patator           794
      SSH-Patator           590
      DoS slowloris        580
      DoS Slowhttptest     550
      Bot                  196
      Web Attack - Brute Force 150
      Web Attack - XSS       65
      Infiltration           3
      Web Attack - Sql Injection 2
      Heartbleed             1
      Name: count, dtype: int64
```

Statistics for the training set

```
[14]: benign_percentage = len(Y_train.label[Y_train["label"]=="BENIGN"])/len(Y_train)
      print('Percentage of benign samples: %.4f' % benign_percentage)
      print(Y_train.is_attack.value_counts())
```

Percentage of benign samples: 0.5842

```
is_attack
0.0    2273097
1.0    1618030
Name: count, dtype: int64
```

4 3. Machine Learning Classifiers

In this section, different machine learning models are applied to classify network traffic. The metrics used to evaluate the models are primarily accuracy, precision, recall and the F1 score.

```
[15]: scaler = StandardScaler()
      scaler.fit(X_train)
```

```
[15]: StandardScaler()
```

```
[16]: def plot_confusion_matrix(model_name, Y_true, Y_pred, labels=["Benign",
      ↪ "Attack"]):
      matrix = confusion_matrix(Y_true.is_attack, Y_pred)
      plt.figure(figsize=(8, 6))
      sns.heatmap(matrix, annot=True, cmap='Blues', fmt='d', xticklabels=labels,
      ↪ yticklabels=labels)
      plt.xlabel('Predicted')
      plt.ylabel('True')
      plt.title(f'Confusion Matrix for {model_name}')
      plt.show()

def metrics_report(dataset_type, y_true, y_predict, print_avg=True):
    print(f"Classification Report ({dataset_type}):")
    print(classification_report(y_true, y_predict, digits=4))
    if print_avg:
        print(f"Avg Precision Score: {average_precision_score(y_true,
      ↪ y_predict, average='weighted'))}")
    print("Accuracy:", accuracy_score(y_true, y_predict))
    res = classification_report(y_true, y_predict, digits=4, output_dict = True)
    res["accuracy"] = accuracy_score(y_true, y_predict)
    return res
```

```
[17]: performance_models = {}
```

Functions for saving and loading modules

```
[18]: import joblib

def save_model(model, model_name):
    file_path = f'models/{model_name}.pkl'
    joblib.dump(model, file_path)
    print(f'Model saved to {file_path}')
```



```
def load_model(model_name):
    file_path = f'models/{model_name}.pkl'
    model = joblib.load(file_path)
    print(f'Model loaded from {file_path}')
    return model

os.makedirs('models', exist_ok=True)
```

4.1 3.1 Decision Trees/Ensembles Algorithms

4.1.1 3.1.1 Regression Forest

Binary Classification

```
[19]: rf_model_binary = RandomForestClassifier(verbose=1, n_jobs=-1,
    ↪class_weight='balanced', criterion="entropy", n_estimators = 300, bootstrap
    ↪= True, max_features=None)
rf_model_binary.fit(scaler.transform(X_train), Y_train.is_attack)
```

[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 16 concurrent workers.

[Parallel(n_jobs=-1)]: Done 18 tasks | elapsed: 4.8min

[Parallel(n_jobs=-1)]: Done 168 tasks | elapsed: 27.4min

[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed: 46.3min finished

```
[19]: RandomForestClassifier(class_weight='balanced', criterion='entropy',
    max_features=None, n_estimators=300, n_jobs=-1,
    verbose=1)
```

```
[20]: print("Evaluation Set Performance")
metrics_report("Evaluation", Y_eval.is_attack, rf_model_binary.predict(scaler.
    ↪transform(X_eval)))
# Predict and evaluate on the test set
print("Test Set Performance")
Y_pred = rf_model_binary.predict(scaler.transform(X_test))
performance_models["rf"] = metrics_report("Test", Y_test.is_attack, Y_pred)
plot_confusion_matrix("Regression Forest", Y_test, Y_pred)
```

Evaluation Set Performance

[Parallel(n_jobs=16)]: Using backend ThreadingBackend with 16 concurrent workers.

[Parallel(n_jobs=16)]: Done 18 tasks | elapsed: 0.0s

[Parallel(n_jobs=16)]: Done 168 tasks | elapsed: 0.3s

[Parallel(n_jobs=16)]: Done 300 out of 300 | elapsed: 0.7s finished

Classification Report (Evaluation):

	precision	recall	f1-score	support
0	1.0000	0.9995	0.9997	227310
1	0.9979	0.9999	0.9989	55764

accuracy			0.9996	283074
macro avg	0.9989	0.9997	0.9993	283074
weighted avg	0.9996	0.9996	0.9996	283074

Avg Precision Score: 0.9978486651284253

Accuracy: 0.9995725499339395

Test Set Performance

[Parallel(n_jobs=16)]: Using backend ThreadingBackend with 16 concurrent workers.

[Parallel(n_jobs=16)]: Done 18 tasks | elapsed: 0.0s

[Parallel(n_jobs=16)]: Done 168 tasks | elapsed: 0.4s

[Parallel(n_jobs=16)]: Done 300 out of 300 | elapsed: 0.7s finished


Classification Report (Test):

	precision	recall	f1-score	support
0	1.0000	0.9994	0.9997	227310
1	0.9976	0.9999	0.9988	55765

accuracy			0.9995	283075
macro avg	0.9988	0.9997	0.9992	283075
weighted avg	0.9995	0.9995	0.9995	283075

Avg Precision Score: 0.9975953147083236

Accuracy: 0.9995230945862404



```
[21]: save_model(rf_model_binary, 'random_forest')
```

Model saved to models/random_forest.pkl

Multi-class classifier

```
[22]: rf_model_multiclass = RandomForestClassifier(verbose=1, n_jobs=-1,
    ↪class_weight='balanced')
rf_model_multiclass.fit(scaler.transform(X_train), Y_train.label_code)
# Predict and evaluate on the evaluation set
print("Evaluation Set Performance")
metrics_report("Evaluation", Y_eval.label_code, rf_model_multiclass.
    ↪predict(scaler.transform(X_eval)), print_avg=False)
# Predict and evaluate on the test set
print("Test Set Performance")
metrics_report("Test", Y_test.label_code, rf_model_multiclass.predict(scaler.
    ↪transform(X_test)), print_avg=False)
```

[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 16 concurrent workers.

```
[Parallel(n_jobs=-1)]: Done 18 tasks      | elapsed: 1.0min
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed: 3.5min finished
[Parallel(n_jobs=16)]: Using backend ThreadingBackend with 16 concurrent
workers.
```

Evaluation Set Performance

```
[Parallel(n_jobs=16)]: Done 18 tasks      | elapsed: 0.2s
[Parallel(n_jobs=16)]: Done 100 out of 100 | elapsed: 1.0s finished
```

Classification Report (Evaluation):

	precision	recall	f1-score	support
0	1.0000	0.9977	0.9988	227310
1	0.3803	1.0000	0.5510	197
2	0.9999	0.9998	0.9999	12803
3	0.9990	1.0000	0.9995	1029
4	0.9966	1.0000	0.9983	23107
5	1.0000	1.0000	1.0000	550
6	1.0000	0.9983	0.9991	579
7	1.0000	1.0000	1.0000	794
8	1.0000	1.0000	1.0000	1
9	1.0000	1.0000	1.0000	4
10	0.9943	0.9999	0.9971	15893
11	1.0000	1.0000	1.0000	589
12	0.7254	0.6821	0.7031	151
13	0.6667	1.0000	0.8000	2
14	0.5625	0.9692	0.7119	65
accuracy				0.9979 283074
macro avg	0.8883	0.9765	0.9173	283074
weighted avg	0.9987	0.9979	0.9982	283074

Accuracy: 0.9979404678635269

Test Set Performance

```
[Parallel(n_jobs=16)]: Using backend ThreadingBackend with 16 concurrent
workers.
```

```
[Parallel(n_jobs=16)]: Done 18 tasks      | elapsed: 0.2s
[Parallel(n_jobs=16)]: Done 100 out of 100 | elapsed: 1.0s finished
```

Classification Report (Test):

	precision	recall	f1-score	support
0	1.0000	0.9975	0.9988	227310
1	0.3748	1.0000	0.5452	196
2	0.9999	0.9999	0.9999	12803
3	0.9990	1.0000	0.9995	1030
4	0.9959	1.0000	0.9980	23108
5	0.9964	0.9982	0.9973	550

6	0.9983	1.0000	0.9991	580
7	1.0000	1.0000	1.0000	794
8	1.0000	1.0000	1.0000	1
9	1.0000	1.0000	1.0000	3
10	0.9931	0.9999	0.9965	15893
11	1.0000	1.0000	1.0000	590
12	0.8015	0.7000	0.7473	150
13	1.0000	1.0000	1.0000	2
14	0.5676	0.9692	0.7159	65
accuracy			0.9978	283075
macro avg	0.9151	0.9777	0.9332	283075
weighted avg	0.9986	0.9978	0.9981	283075

Accuracy: 0.9978380287909565

```
[22]: {'0': {'precision': 0.9999955898372209,
  'recall': 0.9975276054727025,
  'f1-score': 0.9987600730301569,
  'support': 227310.0},
  '1': {'precision': 0.37476099426386233,
  'recall': 1.0,
  'f1-score': 0.545201668984701,
  'support': 196.0},
  '2': {'precision': 0.9999218933062564,
  'recall': 0.9999218933062564,
  'f1-score': 0.9999218933062564,
  'support': 12803.0},
  '3': {'precision': 0.9990300678952473,
  'recall': 1.0,
  'f1-score': 0.9995147986414362,
  'support': 1030.0},
  '4': {'precision': 0.9959486251185242,
  'recall': 1.0,
  'f1-score': 0.9979702008205571,
  'support': 23108.0},
  '5': {'precision': 0.9963702359346642,
  'recall': 0.9981818181818182,
  'f1-score': 0.997275204359673,
  'support': 550.0},
  '6': {'precision': 0.9982788296041308,
  'recall': 1.0,
  'f1-score': 0.9991386735572783,
  'support': 580.0},
  '7': {'precision': 1.0, 'recall': 1.0, 'f1-score': 1.0, 'support': 794.0},
  '8': {'precision': 1.0, 'recall': 1.0, 'f1-score': 1.0, 'support': 1.0},
  '9': {'precision': 1.0, 'recall': 1.0, 'f1-score': 1.0, 'support': 3.0},
```

```

'10': {'precision': 0.9930638005373993,
'recall': 0.9999370792172655,
'f1-score': 0.9964885879107098,
'support': 15893.0},
'11': {'precision': 1.0, 'recall': 1.0, 'f1-score': 1.0, 'support': 590.0},
'12': {'precision': 0.8015267175572519,
'recall': 0.7,
'f1-score': 0.7473309608540926,
'support': 150.0},
'13': {'precision': 1.0, 'recall': 1.0, 'f1-score': 1.0, 'support': 2.0},
'14': {'precision': 0.5675675675675675,
'recall': 0.9692307692307692,
'f1-score': 0.7159090909090909,
'support': 65.0},
'accuracy': 0.9978380287909565,
'macro avg': {'precision': 0.915097621441475,
'recall': 0.9776532776939207,
'f1-score': 0.9331674101582635,
'support': 283075.0},
'weighted avg': {'precision': 0.9986212901610746,
'recall': 0.9978380287909565,
'f1-score': 0.9981151144627087,
'support': 283075.0}}

```

Given that the dataset lacks many samples for some of the attack types, the multiclass classifier has low accuracy for several attack types, even after upsampling, which will cause overfitting in this case. It is better to test other machine learning algorithms for binary classifiers.

4.1.2 3.1.2 Gradient Boost (XGB)

```

[23]: import xgboost as xgb
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import make_scorer, precision_score

# Define the parameter grid for RandomizedSearchCV
param_grid = {
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'n_estimators': [100, 200, 300],
    'max_depth': [3, 5, 7, 10],
    'min_child_weight': [1, 3, 5],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0],
    'gamma': [0, 0.1, 0.2, 0.3],
    'reg_alpha': [0, 0.1, 0.5, 1],
    'reg_lambda': [0, 0.1, 0.5, 1]
}

```

```

# Initialize the XGBoost classifier
xgb_model_binary = xgb.XGBClassifier(objective='binary:logistic',
    ↪scale_pos_weight=1, n_jobs=-1)

# Define the scorer
scorer = make_scorer(precision_score)

# Initialize RandomizedSearchCV
random_search = RandomizedSearchCV(estimator=xgb_model_binary,
    ↪param_distributions=param_grid, n_iter=20,
    ↪scoring=scorer, n_jobs=-1, cv=2, verbose=2,
    ↪random_state=42)

# Fit the model on the training data
random_search.fit(scaler.transform(X_train.astype('float32')), Y_train.
    ↪is_attack)

# Get the best model
best_xgb_model = random_search.best_estimator_

# Print the best hyperparameters
print("Best Hyperparameters:", random_search.best_params_)

```

Fitting 2 folds for each of 20 candidates, totalling 40 fits

Best Hyperparameters: {'subsample': 1.0, 'reg_lambda': 0.1, 'reg_alpha': 0.5, 'n_estimators': 300, 'min_child_weight': 1, 'max_depth': 7, 'learning_rate': 0.05, 'gamma': 0.1, 'colsample_bytree': 1.0}

Best Hyperparameters: {'subsample': 1.0, 'reg_lambda': 0, 'reg_alpha': 0.5, 'n_estimators': 100, 'min_child_weight': 3, 'max_depth': 10, 'learning_rate': 0.1, 'gamma': 0.3, 'colsample_bytree': 0.6}

```

[24]: best_params = random_search.best_params_

best_xgb_model = xgb.XGBClassifier(
    objective='binary:logistic',
    scale_pos_weight=1,
    n_jobs=-1,
    **best_params
)

# Fit the model on the training data
best_xgb_model.fit(scaler.transform(X_train), Y_train.is_attack)

```

```

[24]: XGBClassifier(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=1.0, device=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric=None, feature_types=None,

```

```

gamma=0.1, grow_policy=None, importance_type=None,
interaction_constraints=None, learning_rate=0.05, max_bin=None,
max_cat_threshold=None, max_cat_to_onehot=None,
max_delta_step=None, max_depth=7, max_leaves=None,
min_child_weight=1, missing=nan, monotone_constraints=None,
multi_strategy=None, n_estimators=300, n_jobs=-1,
num_parallel_tree=None, random_state=None, ...)

```

```

[25]: # Predict and evaluate on the evaluation set
print("Evaluation Set Performance")
metrics_report("Evaluation", Y_eval.is_attack, best_xgb_model.predict(scaler.
    ↪transform(X_eval)), print_avg=False)
# Predict and evaluate on the test set
print("Test Set Performance")
Y_pred = best_xgb_model.predict(scaler.transform(X_test))
performance_models["xgb"] = metrics_report("Test", Y_test.is_attack, Y_pred, ↪
    ↪print_avg=False)
plot_confusion_matrix("Gradient Boost", Y_test, Y_pred)

```

Evaluation Set Performance

Classification Report (Evaluation):

	precision	recall	f1-score	support
0	0.9998	0.9983	0.9990	227310
1	0.9930	0.9991	0.9960	55764
accuracy			0.9984	283074
macro avg	0.9964	0.9987	0.9975	283074
weighted avg	0.9984	0.9984	0.9984	283074


Accuracy: 0.9984279728975463

Test Set Performance

Classification Report (Test):

	precision	recall	f1-score	support
0	0.9998	0.9982	0.9990	227310
1	0.9927	0.9991	0.9959	55765
accuracy			0.9984	283075
macro avg	0.9962	0.9987	0.9974	283075
weighted avg	0.9984	0.9984	0.9984	283075

Accuracy: 0.9983785215932174



```
[26]: save_model(best_xgb_model, 'xgb_model')
```

Model saved to models/xgb_model.pkl

4.1.3 3.1.3 ADABOOST

```
[27]: from sklearn.ensemble import AdaBoostClassifier
ada_boost_model = AdaBoostClassifier(
    n_estimators=50, # Number of weak learners
    learning_rate=1.0, # Learning rate (contribution of each weak learner)
    algorithm='SAMME', # SAMME.R is recommended for probability estimates
    random_state=42
)

# Fit the model on the training data
ada_boost_model.fit(scaler.transform(X_train), Y_train.is_attack)
```

```
[27]: AdaBoostClassifier(algorithm='SAMME', random_state=42)
```

```
[28]: # Predict and evaluate on the evaluation set
y_pred_eval = ada_boost_model.predict(scaler.transform(X_eval))
metrics_report("Evaluation", Y_eval.is_attack, y_pred_eval, print_avg=False)

# Predict and evaluate on the test set
y_pred_test = ada_boost_model.predict(scaler.transform(X_test))
performance_models["adaboost"] = metrics_report("Test", Y_test.is_attack,
↪y_pred_test, print_avg=False)
plot_confusion_matrix("Ada Boost", Y_test, y_pred_test)
```

Classification Report (Evaluation):


	precision	recall	f1-score	support
0	0.9910	0.9221	0.9553	227310
1	0.7525	0.9657	0.8459	55764
accuracy			0.9307	283074
macro avg	0.8718	0.9439	0.9006	283074
weighted avg	0.9440	0.9307	0.9337	283074

Accuracy: 0.9306930343302457

Classification Report (Test):

	precision	recall	f1-score	support
0	0.9915	0.9225	0.9557	227310
1	0.7539	0.9677	0.8475	55765
accuracy			0.9314	283075
macro avg	0.8727	0.9451	0.9016	283075
weighted avg	0.9447	0.9314	0.9344	283075

Accuracy: 0.9313927404398128



```
[29]: save_model(ada_boost_model, 'ada_boost_model')
```

Model saved to models/ada_boost_model.pkl

4.1.4 3.1.4 ID3

```
[30]: id3_model = DecisionTreeClassifier(
    criterion='entropy', # Use entropy for ID3-like behavior
    max_depth=None, # Unlimited depth (ID3 doesn't prune trees)
    random_state=42
)

# Fit the model on the training data
id3_model.fit(scaler.transform(X_train), Y_train.is_attack)
```

```
[30]: DecisionTreeClassifier(criterion='entropy', random_state=42)
```

```
[31]: y_pred_eval = id3_model.predict(scaler.transform(X_eval))
metrics_report("Evaluation", Y_eval.is_attack, y_pred_eval, print_avg=False)
```

```
# Predict and evaluate on the test set
y_pred_test = id3_model.predict(scaler.transform(X_test))
performance_models["id3"] = metrics_report("Test", Y_test.is_attack,
    ↪y_pred_test, print_avg=False)
plot_confusion_matrix("ID3", Y_test, y_pred_test)
```

Classification Report (Evaluation):


	precision	recall	f1-score	support
0	1.0000	0.9995	0.9997	227310
1	0.9980	0.9999	0.9989	55764
accuracy			0.9996	283074
macro avg	0.9990	0.9997	0.9993	283074
weighted avg	0.9996	0.9996	0.9996	283074

Accuracy: 0.9995760825791136

Classification Report (Test):

	precision	recall	f1-score	support
0	1.0000	0.9994	0.9997	227310
1	0.9977	0.9999	0.9988	55765
accuracy			0.9995	283075
macro avg	0.9988	0.9997	0.9992	283075
weighted avg	0.9995	0.9995	0.9995	283075

Accuracy: 0.9995230945862404



```
[32]: save_model(id3_model, 'id3_model')
```

Model saved to models/id3_model.pkl

4.1.5 Conclusion

```
[33]: def extract_and_plot_metrics(metrics_dict):
    # Initialize dictionaries to store the metrics for plotting
    precision_dict = {'0': [], '1': [], 'model': []}
    recall_dict = {'0': [], '1': [], 'model': []}
    f1_score_dict = {'0': [], '1': [], 'model': []}
    accuracy_list = []

    # Iterate over the models in the metrics dictionary
    for model_name, metrics in metrics_dict.items():
        precision_dict['0'].append(metrics['0']['precision'])
        precision_dict['1'].append(metrics['1']['precision'])
        recall_dict['0'].append(metrics['0']['recall'])
        recall_dict['1'].append(metrics['1']['recall'])
```

```

f1_score_dict['0'].append(metrics['0']['f1-score'])
f1_score_dict['1'].append(metrics['1']['f1-score'])
accuracy_list.append(metrics['accuracy'])
precision_dict['model'].append(model_name)
recall_dict['model'].append(model_name)
f1_score_dict['model'].append(model_name)

# Plotting the metrics
fig, axs = plt.subplots(2, 2, figsize=(14, 10))

# Plot precision
axs[0, 0].plot(precision_dict['model'], precision_dict['0'], label='Class 0',
↪marker='o')
axs[0, 0].plot(precision_dict['model'], precision_dict['1'], label='Class 1',
↪marker='o')
axs[0, 0].set_title('Precision')
axs[0, 0].legend()

# Plot recall
axs[0, 1].plot(recall_dict['model'], recall_dict['0'], label='Class 0',
↪marker='o')
axs[0, 1].plot(recall_dict['model'], recall_dict['1'], label='Class 1',
↪marker='o')
axs[0, 1].set_title('Recall')
axs[0, 1].legend()

# Plot f1-score
axs[1, 0].plot(f1_score_dict['model'], f1_score_dict['0'], label='Class 0',
↪marker='o')
axs[1, 0].plot(f1_score_dict['model'], f1_score_dict['1'], label='Class 1',
↪marker='o')
axs[1, 0].set_title('F1-Score')
axs[1, 0].legend()


# Plot accuracy
print(accuracy_list)
print(precision_dict['model'])
axs[1, 1].plot(precision_dict['model'], accuracy_list, label='Accuracy',
↪marker='o')
axs[1, 1].set_title('Accuracy')
axs[1, 1].legend()

plt.tight_layout()
plt.show()

extract_and_plot_metrics(performance_models)

```

[0.9995230945862404, 0.9983785215932174, 0.9313927404398128, 0.9995230945862404]
['rf', 'xgb', 'adaboost', 'id3']



ID3 and Random Forest perform very well on the dataset while Adaboost is the worst model. To conclude, after finetuning the hyperparameters, random forest seems to be the best tree based algorithm to perform on the dataset.

4.2 3.2 Deep Neural Network

In this section, deep neural networks are used to create binary classifiers that distinguish between benign and malicious traffics. Hyperparameters are optimized to obtain the best results.

```
[34]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
```

```
[35]: # Define the model architecture
model = keras.Sequential([
    keras.layers.Dense(128, activation='relu', input_shape=(scaler.
    ↪transform(X_train).shape[1],)),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(64, activation='relu'),
```

```

        keras.layers.Dropout(0.5),
        keras.layers.Dense(1, activation='sigmoid')
    ])

    # Compile the model
    model.compile(optimizer='adam', loss='binary_crossentropy',
        metrics=['accuracy'])

    # Train the model
    history = model.fit(scaler.transform(X_train), Y_train.is_attack, epochs=10,
        batch_size=32, validation_split=0.2)

```

C:\Users\youuss\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\layers\core\dense.py:88: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

```

Epoch 1/10
97279/97279          66s
673us/step - accuracy: 0.9374 - loss: 0.1676 - val_accuracy: 0.2848 - val_loss: 8.9642
Epoch 2/10
97279/97279          66s
678us/step - accuracy: 0.9524 - loss: 0.1146 - val_accuracy: 0.2882 - val_loss: 7.4969
Epoch 3/10
97279/97279          65s
669us/step - accuracy: 0.9544 - loss: 0.1118 - val_accuracy: 0.2889 - val_loss: 4.7225
Epoch 4/10
97279/97279          65s
664us/step - accuracy: 0.9557 - loss: 0.1169 - val_accuracy: 0.2726 - val_loss: 10.9780
Epoch 5/10
97279/97279          64s
661us/step - accuracy: 0.9562 - loss: 0.1089 - val_accuracy: 0.3051 - val_loss: 6.2204
Epoch 6/10
97279/97279          65s
666us/step - accuracy: 0.9568 - loss: 0.1259 - val_accuracy: 0.3020 - val_loss: 9.3219
Epoch 7/10
97279/97279          65s
666us/step - accuracy: 0.9570 - loss: 0.1114 - val_accuracy: 0.2990 - val_loss: 11.0146
Epoch 8/10
97279/97279          65s

```


666us/step - accuracy: 0.9575 - loss: 0.1486 - val_accuracy: 0.3099 - val_loss: 5.7396

Epoch 9/10

97279/97279 65s

665us/step - accuracy: 0.9577 - loss: 0.1267 - val_accuracy: 0.3003 - val_loss: 9.3866

Epoch 10/10

97279/97279 65s

666us/step - accuracy: 0.9578 - loss: 0.1226 - val_accuracy: 0.3451 - val_loss: 8.2815

```
[36]: # Predict probabilities on the evaluation set
y_pred_eval_prob = model.predict(scaler.transform(X_eval))
# Convert probabilities to binary predictions
y_pred_eval = (y_pred_eval_prob > 0.5).astype(int)
metrics_report("Evaluation", Y_eval.is_attack, y_pred_eval, print_avg=False)

# Predict and evaluate on the test set
y_pred_test_prob = model.predict(scaler.transform(X_test))
y_pred_test = (y_pred_test_prob > 0.5).astype(int)

metrics_report("Test", Y_test.is_attack, y_pred_test, print_avg=False)
plot_confusion_matrix("DNN", Y_test, y_pred_test)
```

8847/8847 4s 420us/step

Classification Report (Evaluation):

	precision	recall	f1-score	support
0	0.9813	0.9796	0.9805	227310
1	0.9175	0.9239	0.9207	55764
accuracy			0.9686	283074
macro avg	0.9494	0.9518	0.9506	283074
weighted avg	0.9687	0.9686	0.9687	283074


Accuracy: 0.9686477740802758

8847/8847 4s 433us/step

Classification Report (Test):

	precision	recall	f1-score	support
0	0.9818	0.9798	0.9808	227310
1	0.9185	0.9260	0.9222	55765
accuracy			0.9692	283075
macro avg	0.9501	0.9529	0.9515	283075
weighted avg	0.9693	0.9692	0.9693	283075

Accuracy: 0.9692307692307692



4.2.1 Hyperparameter Tuning

```
[37]: def build_model(hp):
    model = Sequential()
    model.add(Dense(units=hp.Int('units_input', min_value=32, max_value=512,
    ↪step=32), activation='relu', input_shape=(X_train.shape[1],)))

    for i in range(hp.Int('num_layers', 1, 3)):
        model.add(Dense(units=hp.Int(f'units_{i}', min_value=32, max_value=512,
    ↪step=32), activation='relu'))
        model.add(Dropout(rate=hp.Float(f'dropout_{i}', min_value=0.0,
    ↪max_value=0.5, step=0.1)))

    model.add(Dense(1, activation='sigmoid'))

    model.compile(
```

```

        optimizer=tf.keras.optimizers.Adam(learning_rate=hp.
↪Float('learning_rate', min_value=1e-4, max_value=1e-2, sampling='LOG',
↪default=1e-3)),
        loss='binary_crossentropy',
        metrics=['accuracy']
    )

    return model

```

```

[52]: from keras_tuner import RandomSearch

directory_path = os.path.join(os.getcwd(), 'hyperparam_tuning')
tuner = RandomSearch(
    build_model,
    objective='val_accuracy',
    max_trials=10,
    executions_per_trial=2,
    directory=directory_path,
    project_name='intrusion_detection'
)

tuner.search_space_summary()

# Perform the search
tuner.search(scaler.transform(X_train), Y_train.is_attack, epochs=10,
↪validation_split=0.2)

```

```

Reloading Tuner from G:\Other computers\My PC\stage\ML-
NIDS\Notebooks\hyperparam_tuning\intrusion_detection\tuner0.json
Search space summary
Default search space size: 9
units_input (Int)
{'default': None, 'conditions': [], 'min_value': 32, 'max_value': 512, 'step':
32, 'sampling': 'linear'}
num_layers (Int)
{'default': None, 'conditions': [], 'min_value': 1, 'max_value': 3, 'step': 1,
'sampling': 'linear'}
units_0 (Int)
{'default': None, 'conditions': [], 'min_value': 32, 'max_value': 512, 'step':
32, 'sampling': 'linear'}
dropout_0 (Float)
{'default': 0.0, 'conditions': [], 'min_value': 0.0, 'max_value': 0.5, 'step':
0.1, 'sampling': 'linear'}
learning_rate (Float)
{'default': 0.001, 'conditions': [], 'min_value': 0.0001, 'max_value': 0.01,
'step': None, 'sampling': 'log'}
units_1 (Int)
{'default': None, 'conditions': [], 'min_value': 32, 'max_value': 512, 'step':

```

```

32, 'sampling': 'linear'}
dropout_1 (Float)
{'default': 0.0, 'conditions': [], 'min_value': 0.0, 'max_value': 0.5, 'step':
0.1, 'sampling': 'linear'}
units_2 (Int)
{'default': None, 'conditions': [], 'min_value': 32, 'max_value': 512, 'step':
32, 'sampling': 'linear'}
dropout_2 (Float)
{'default': 0.0, 'conditions': [], 'min_value': 0.0, 'max_value': 0.5, 'step':
0.1, 'sampling': 'linear'}

```

```

[39]: best_hps = tuner.get_best_hyperparameters(num_trials=1)[0]
print(best_hps)
model1 = build_model(best_hps)
history = model.fit(scaler.transform(X_train), Y_train.is_attack, epochs=20,
validation_split=0.2, verbose=1)

```

```

<keras_tuner.src.engine.hyperparameters.hyperparameters.HyperParameters object
at 0x0000015161FA3830>

```

```

C:\Users\youuss\AppData\Local\Programs\Python\Python312\Lib\site-
packages\keras\src\layers\core\dense.py:88: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

```

Epoch 1/20
97279/97279          65s
667us/step - accuracy: 0.9579 - loss: 0.1334 - val_accuracy: 0.3554 - val_loss:
5.8488
Epoch 2/20
97279/97279          65s
668us/step - accuracy: 0.9580 - loss: 0.1261 - val_accuracy: 0.3384 - val_loss:
7.1888
Epoch 3/20
97279/97279          65s
670us/step - accuracy: 0.9584 - loss: 0.1040 - val_accuracy: 0.3418 - val_loss:
4.9406
Epoch 4/20
97279/97279          65s
668us/step - accuracy: 0.9583 - loss: 0.1237 - val_accuracy: 0.3384 - val_loss:
5.9503
Epoch 5/20
97279/97279          65s
668us/step - accuracy: 0.9586 - loss: 0.1698 - val_accuracy: 0.3462 - val_loss:
6.6066
Epoch 6/20
97279/97279          65s
666us/step - accuracy: 0.9587 - loss: 0.1062 - val_accuracy: 0.3664 - val_loss:

```

6.7357
Epoch 7/20
97279/97279 65s
668us/step - accuracy: 0.9586 - loss: 0.1191 - val_accuracy: 0.3445 - val_loss: 6.9325
Epoch 8/20
97279/97279 65s
667us/step - accuracy: 0.9586 - loss: 0.1078 - val_accuracy: 0.3645 - val_loss: 5.8693
Epoch 9/20
97279/97279 65s
667us/step - accuracy: 0.9586 - loss: 0.2841 - val_accuracy: 0.3554 - val_loss: 6.4449
Epoch 10/20
97279/97279 65s
665us/step - accuracy: 0.9589 - loss: 0.1272 - val_accuracy: 0.3623 - val_loss: 7.7536
Epoch 11/20
97279/97279 65s
665us/step - accuracy: 0.9591 - loss: 0.1058 - val_accuracy: 0.3481 - val_loss: 7.2022
Epoch 12/20
97279/97279 65s
671us/step - accuracy: 0.9593 - loss: 0.1060 - val_accuracy: 0.3884 - val_loss: 7.3497
Epoch 13/20
97279/97279 66s
675us/step - accuracy: 0.9587 - loss: 0.1452 - val_accuracy: 0.3858 - val_loss: 5.5309
Epoch 14/20
97279/97279 66s
676us/step - accuracy: 0.9594 - loss: 0.1067 - val_accuracy: 0.4324 - val_loss: 4.6448
Epoch 15/20
97279/97279 65s
667us/step - accuracy: 0.9589 - loss: 0.1081 - val_accuracy: 0.4341 - val_loss: 3.4974
Epoch 16/20
97279/97279 65s
667us/step - accuracy: 0.9595 - loss: 0.2097 - val_accuracy: 0.4070 - val_loss: 4.3750
Epoch 17/20
97279/97279 66s
673us/step - accuracy: 0.9593 - loss: 0.1092 - val_accuracy: 0.3577 - val_loss: 6.5637
Epoch 18/20
97279/97279 65s
665us/step - accuracy: 0.9596 - loss: 0.1057 - val_accuracy: 0.3408 - val_loss:

```

7.4742
Epoch 19/20
97279/97279          65s
669us/step - accuracy: 0.9595 - loss: 0.1479 - val_accuracy: 0.3792 - val_loss:
4.0979
Epoch 20/20
97279/97279          65s
668us/step - accuracy: 0.9592 - loss: 0.1318 - val_accuracy: 0.4012 - val_loss:
7.1979

```

```
[40]: print(best_hps.values)
```

```

{'units_input': 480, 'num_layers': 2, 'units_0': 448, 'dropout_0':
0.30000000000000004, 'learning_rate': 0.00614260757976685, 'units_1': 224,
'dropout_1': 0.0, 'units_2': 256, 'dropout_2': 0.4}

```

```
[41]: from tensorflow.keras.models import save_model as save_model_keras
```

```

def save_keras_model(model, model_name):
    file_path = f'models/{model_name}.h5'
    save_model_keras(model, file_path)
    print(f'Model saved to {file_path}')

save_keras_model(model, 'DNN_model1')

```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

Model saved to models/DNN_model1.h5

Best Hyperparameters {'units_input': 480, 'num_layers': 2, 'units_0': 448, 'dropout_0': 0.30000000000000004, 'learning_rate': 0.00614260757976685, 'units_1': 224, 'dropout_1': 0.0, 'units_2': 256, 'dropout_2': 0.4}

```

[42]: # Predict probabilities on the evaluation set
y_pred_eval_prob = model.predict(scaler.transform(X_eval))
# Convert probabilities to binary predictions
y_pred_eval = (y_pred_eval_prob > 0.5).astype(int)
metrics_report("Evaluation", Y_eval.is_attack, y_pred_eval, print_avg=False)
# Predict and evaluate on the test set
y_pred_test_prob = model.predict(scaler.transform(X_test))
y_pred_test = (y_pred_test_prob > 0.5).astype(int)

metrics_report("Test", Y_test.is_attack, y_pred_test, print_avg=False)
plot_confusion_matrix("DNN", Y_test, y_pred_test)

```

```

8847/8847          4s 424us/step
Classification Report (Evaluation):
      precision    recall  f1-score   support

     0       0.9833     0.9737     0.9785     227310
     1       0.8969     0.9326     0.9144      55764

 accuracy                   0.9656     283074
 macro avg       0.9401     0.9531     0.9464     283074
weighted avg       0.9663     0.9656     0.9658     283074

```

Accuracy: 0.9655920360047197

```

8847/8847          4s 446us/step
Classification Report (Test):
      precision    recall  f1-score   support

     0       0.9838     0.9738     0.9788     227310
     1       0.8976     0.9348     0.9158      55765

 accuracy                   0.9662     283075
 macro avg       0.9407     0.9543     0.9473     283075
weighted avg       0.9669     0.9662     0.9664     283075

```

Accuracy: 0.9661503135211517

