

5_More_Supervised_Models

June 5, 2024

1 Supervised Learning Algorithms

In this notebook, several supervised learning algorithms are used to create binary classifiers that distinguish between benign and malicious traffics in the ids-2017 dataset. Hyperparameters are optimized to obtain a model with the best results and the models are compared

```
[1]: from notebook_utils import load_processed_dataset_2017, plot_confusion_matrix, \
      ↪ metrics_report, upsample_dataset, extract_and_plot_metrics
      from sklearn.decomposition import PCA

      import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt
      import seaborn as sns
      import glob
      import os

      from sklearn.model_selection import train_test_split, RandomizedSearchCV
      from sklearn.preprocessing import StandardScaler
      from sklearn.metrics import classification_report, average_precision_score, \
      ↪ make_scorer, precision_score, accuracy_score, confusion_matrix
      from notebook_utils import upsample_dataset
      %matplotlib inline
      %load_ext autoreload
      %autoreload 2
      file_path = \
      ↪ r"CIC-IDS-2017\CSVs\GeneratedLabelledFlows\TrafficLabelling\processed\ids2017_processed.
      ↪ csv"

[2]: X_train, Y_train, X_eval, Y_eval, X_test, Y_test, scaler = \
      ↪ load_processed_dataset_2017(file_path)

[3]: performance_models = {}

[4]: import joblib

      def save_model(model, model_name):
          file_path = f'models/{model_name}.pkl'
          joblib.dump(model, file_path)
```

```

    print(f'Model saved to {file_path}')

def load_model(model_name):
    file_path = f'models/{model_name}.pkl'
    model = joblib.load(file_path)
    print(f'Model loaded from {file_path}')
    return model

os.makedirs('models', exist_ok=True)

```

1.1 1. Naive Bayes

```

[5]: from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import GridSearchCV

# Define a small parameter grid
param_grid_nb = {
    'var_smoothing': [1e-9, 1e-8, 1e-7, 1e-6, 1e-5]
}

# Initialize the GaussianNB model
nb_model = GaussianNB()

# Initialize GridSearchCV
grid_search_nb = GridSearchCV(estimator=nb_model, param_grid=param_grid_nb,
    ↪scoring='accuracy', n_jobs=2, cv=2, verbose=2)

# Fit the model on the scaled training data
grid_search_nb.fit(scaler.transform(X_train), Y_train.is_attack)

# Get the best model
best_nb_model = grid_search_nb.best_estimator_
save_model(best_nb_model, 'gaussian_nb_model.pkl')

```

Fitting 2 folds for each of 5 candidates, totalling 10 fits
 Model saved to models/gaussian_nb_model.pkl.pkl

```

[6]: # Predict and evaluate on the evaluation set
print("Evaluation Set Performance")
metrics_report("Evaluation", Y_eval.is_attack, best_nb_model.predict(scaler.
    ↪transform(X_eval)), print_avg=False)

# Predict and evaluate on the test set
print("Test Set Performance")
Y_pred = best_nb_model.predict(scaler.transform(X_test))
performance_models["GaussianNB"] = metrics_report("Test", Y_test.is_attack,
    ↪Y_pred, print_avg=False)

```

```
plot_confusion_matrix("GaussianNB", Y_test, Y_pred)
```

Evaluation Set Performance

Classification Report (Evaluation):

	precision	recall	f1-score	support
0	0.9992	0.4445	0.6153	227310
1	0.3060	0.9985	0.4685	55764
accuracy			0.5537	283074
macro avg	0.6526	0.7215	0.5419	283074
weighted avg	0.8626	0.5537	0.5864	283074

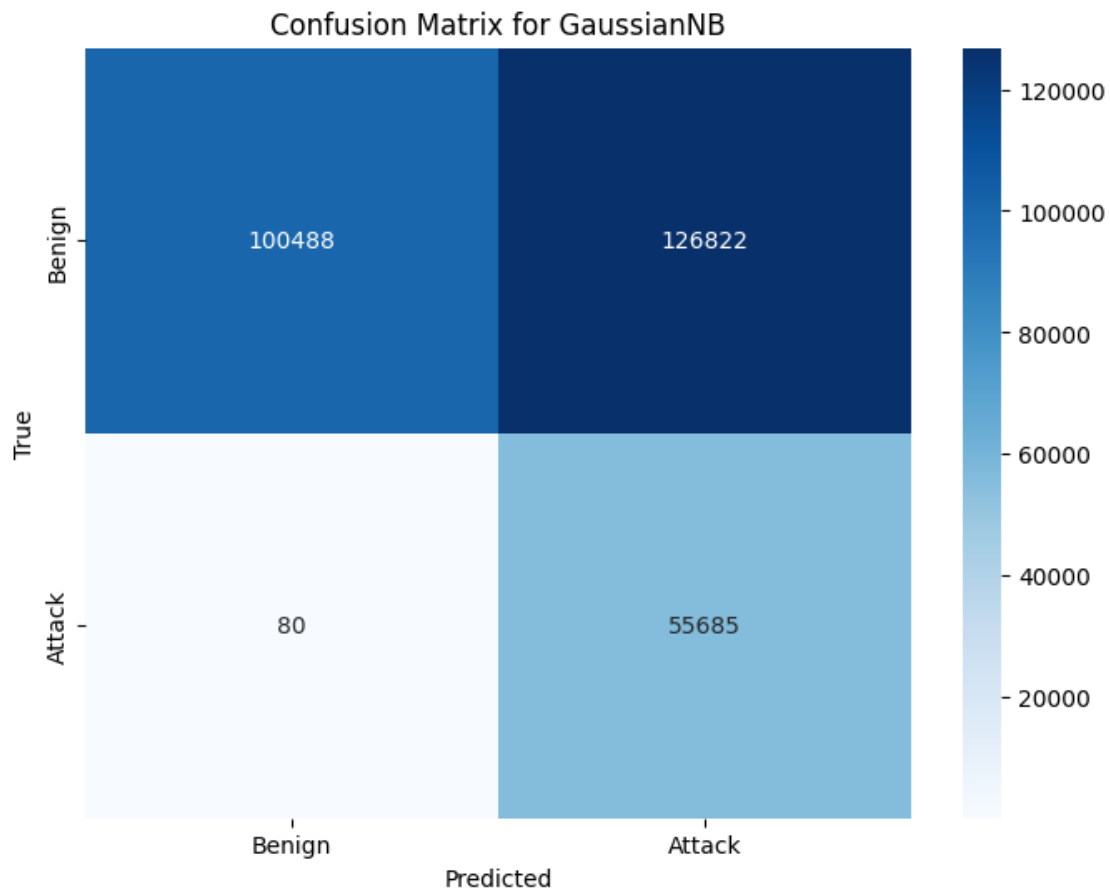
Accuracy: 0.5536714781293938

Test Set Performance

Classification Report (Test):

	precision	recall	f1-score	support
0	0.9992	0.4421	0.6130	227310
1	0.3051	0.9986	0.4674	55765
accuracy			0.5517	283075
macro avg	0.6522	0.7203	0.5402	283075
weighted avg	0.8625	0.5517	0.5843	283075

Accuracy: 0.5517018458005829



```
[15]: performance_models["GaussianNB"] = metrics_report("Test", Y_test.is_attack,
↳ Y_pred, print_avg=False)
```

Classification Report (Test):

	precision	recall	f1-score	support
0	0.9446	0.6808	0.7913	227310
1	0.3916	0.8373	0.5336	55765
accuracy			0.7117	283075
macro avg	0.6681	0.7591	0.6625	283075
weighted avg	0.8357	0.7117	0.7406	283075

Accuracy: 0.7116559215755541

1.2 2. K-Nearest Neighbors (KNN)

```
[8]: from sklearn.neighbors import KNeighborsClassifier
knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(scaler.transform(X_train), Y_train.is_attack)

# Predict and evaluate on the evaluation set
print("Evaluation Set Performance")
metrics_report("Evaluation", Y_eval.is_attack, knn_model.predict(scaler.
    ↪transform(X_eval)), print_avg=False)

# Predict and evaluate on the test set
print("Test Set Performance")
Y_pred = knn_model.predict(scaler.transform(X_test))
performance_models["KNN"] = metrics_report("Test", Y_test.is_attack, Y_pred,
    ↪print_avg=False)
plot_confusion_matrix("KNN", Y_test, Y_pred)

save_model(knn_model, 'knn_model')
```

Evaluation Set Performance

Classification Report (Evaluation):

	precision	recall	f1-score	support
0	0.9735	0.9986	0.9858	227310
1	0.9934	0.8890	0.9383	55764
accuracy			0.9770	283074
macro avg	0.9834	0.9438	0.9621	283074
weighted avg	0.9774	0.9770	0.9765	283074

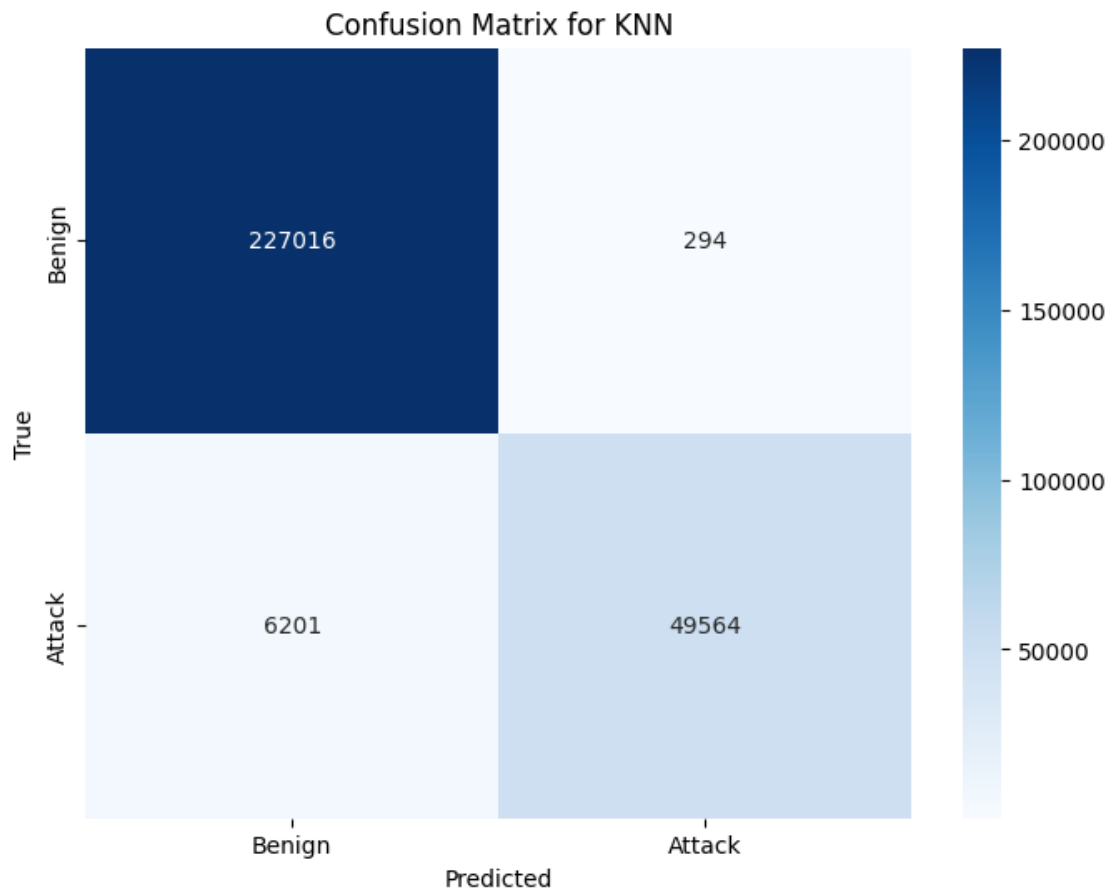
Accuracy: 0.9769742187555197

Test Set Performance

Classification Report (Test):

	precision	recall	f1-score	support
0	0.9734	0.9987	0.9859	227310
1	0.9941	0.8888	0.9385	55765
accuracy			0.9771	283075
macro avg	0.9838	0.9438	0.9622	283075
weighted avg	0.9775	0.9771	0.9766	283075

Accuracy: 0.9770555506491213



Model saved to models/knn_model.pkl

```
[14]: performance_models["KNN"] = metrics_report("Test", Y_test.is_attack, Y_pred,
        print_avg=False)
```

Classification Report (Test):

	precision	recall	f1-score	support
0	0.9446	0.6808	0.7913	227310
1	0.3916	0.8373	0.5336	55765
accuracy			0.7117	283075
macro avg	0.6681	0.7591	0.6625	283075
weighted avg	0.8357	0.7117	0.7406	283075

Accuracy: 0.7116559215755541

1.3 3. Quadratic Discriminant Analysis (QDA)

```
[9]: from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
from sklearn.model_selection import GridSearchCV
# Apply PCA
pca = PCA(n_components=0.90)
X_train_pca = pca.fit_transform(scaler.transform(X_train))
X_eval_pca = pca.transform(scaler.transform(X_eval))
X_test_pca = pca.transform(scaler.transform(X_test))
# Define a small parameter grid
param_grid_qda = {
    'reg_param': [0.0, 0.1, 0.2, 0.3],
    'tol': [1e-4, 1e-5, 1e-6]
}

# Initialize the QDA model
qda_model = QuadraticDiscriminantAnalysis()

# Initialize GridSearchCV with reduced number of folds
grid_search_qda = GridSearchCV(estimator=qda_model, param_grid=param_grid_qda,
    ↪scoring='accuracy', n_jobs=2, cv=2, verbose=2)

# Fit the model on the PCA-transformed training data
grid_search_qda.fit(X_train_pca, Y_train.is_attack)

# Get the best model
best_qda_model = grid_search_qda.best_estimator_

# Print the best hyperparameters
print("Best Hyperparameters for QDA:", grid_search_qda.best_params_)
save_model(best_qda_model, 'qda_model.pkl')
```

Fitting 2 folds for each of 12 candidates, totalling 24 fits
Best Hyperparameters for QDA: {'reg_param': 0.3, 'tol': 0.0001}
Model saved to models/qda_model.pkl.pkl

```
[10]: # Predict and evaluate on the evaluation set
print("Evaluation Set Performance")
metrics_report("Evaluation", Y_eval.is_attack, best_qda_model.
    ↪predict(X_eval_pca), print_avg=False)

# Predict and evaluate on the test set
print("Test Set Performance")
Y_pred = best_qda_model.predict(X_test_pca)
performance_models["QDA"] = metrics_report("Test", Y_test.is_attack, Y_pred,
    ↪print_avg=False)
plot_confusion_matrix("QDA", Y_test, Y_pred)
```

Evaluation Set Performance

Classification Report (Evaluation):

	precision	recall	f1-score	support
0	0.9444	0.6844	0.7936	227310
1	0.3938	0.8357	0.5353	55764
accuracy			0.7142	283074
macro avg	0.6691	0.7600	0.6645	283074
weighted avg	0.8359	0.7142	0.7427	283074

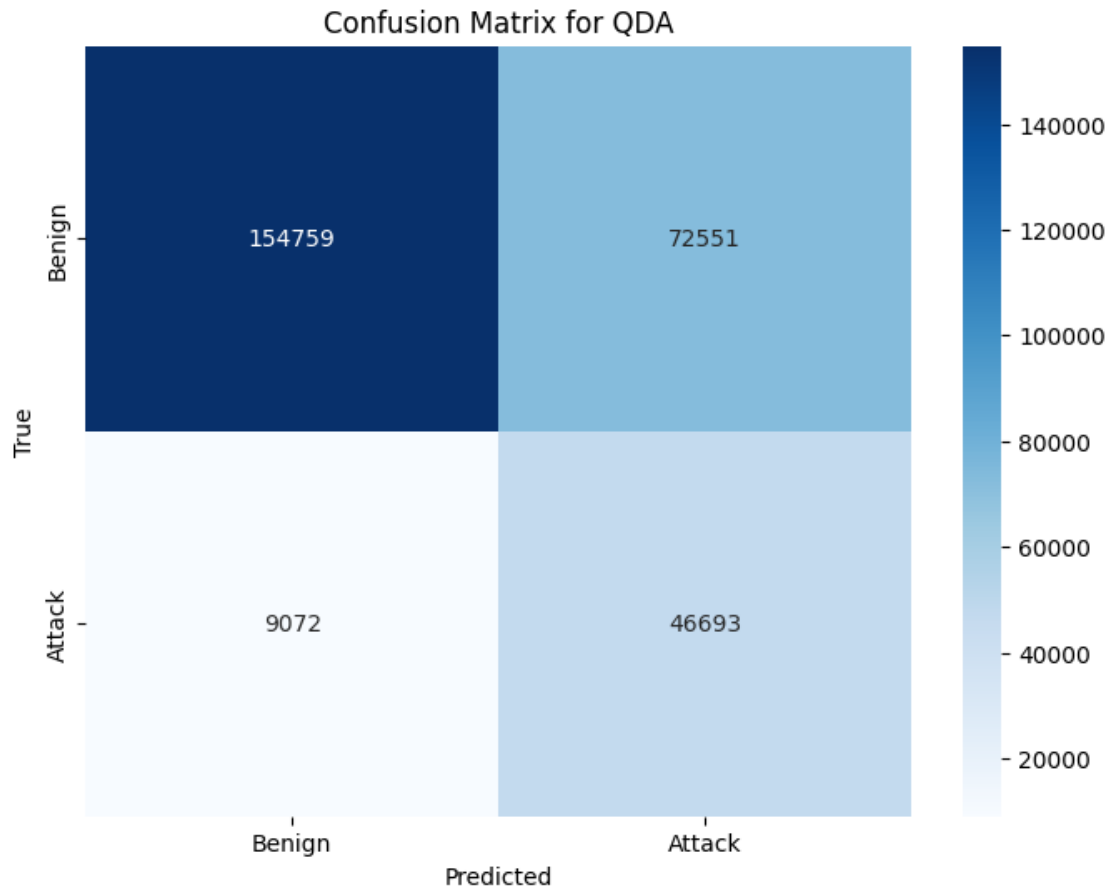
Accuracy: 0.7141772116125112

Test Set Performance

Classification Report (Test):

	precision	recall	f1-score	support
0	0.9446	0.6808	0.7913	227310
1	0.3916	0.8373	0.5336	55765
accuracy			0.7117	283075
macro avg	0.6681	0.7591	0.6625	283075
weighted avg	0.8357	0.7117	0.7406	283075

Accuracy: 0.7116559215755541



1.4 4. Support Vector Machine

The SVM takes too long with the number of features and number of training examples. PCA components are used to reduce the number of features.

```
[11]: from sklearn.svm import LinearSVC
      from sklearn.decomposition import PCA
      import joblib
      import os

      # Apply PCA to reduce the number of features
      pca = PCA(n_components=5)
      X_train_pca = pca.fit_transform(X_train)
      X_eval_pca = pca.transform(X_eval)
      X_test_pca = pca.transform(X_test)

      # Define the LinearSVC model with chosen hyperparameters
      linear_svc_model = LinearSVC(C=1, max_iter=10000, dual="auto")
```

```

# Fit the model on the training data
linear_svc_model.fit(X_train_pca, Y_train.is_attack)

# Save the model
save_model(linear_svc_model, 'linear_svc_binary_pca')

# Evaluate the model
# Predict on the evaluation set
y_pred_eval = linear_svc_model.predict(X_eval_pca)
performance_eval = metrics_report("Evaluation", Y_eval.is_attack, y_pred_eval,
    print_avg=False)

# Predict and evaluate on the test set
y_pred_test = linear_svc_model.predict(X_test_pca)
performance_models["LinearSVC"] = metrics_report("Test", Y_test.is_attack,
    y_pred_test, print_avg=False)

# Plot the confusion matrix
plot_confusion_matrix("LinearSVC", Y_test, y_pred_test)

```

Model saved to models/linear_svc_binary_pca.pkl

Classification Report (Evaluation):


	precision	recall	f1-score	support
0	0.8403	0.8796	0.8595	227310
1	0.3935	0.3186	0.3521	55764
accuracy			0.7690	283074
macro avg	0.6169	0.5991	0.6058	283074
weighted avg	0.7523	0.7690	0.7595	283074

Accuracy: 0.7690497891010831

Classification Report (Test):

	precision	recall	f1-score	support
0	0.8403	0.8781	0.8588	227310
1	0.3915	0.3196	0.3519	55765
accuracy			0.7681	283075
macro avg	0.6159	0.5989	0.6053	283075
weighted avg	0.7519	0.7681	0.7589	283075

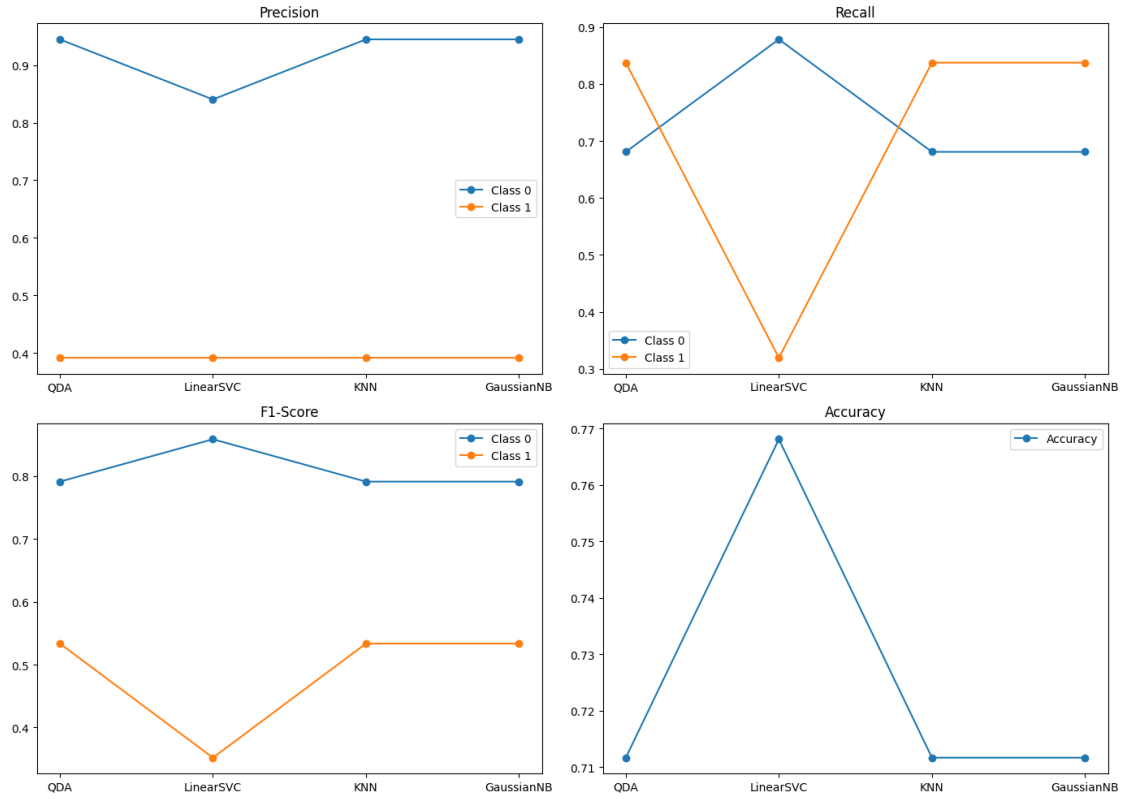
Accuracy: 0.7681109246666078



1.5 5. Conclusion

```
[16]: extract_and_plot_metrics(performance_models)
```

```
[0.7116559215755541, 0.7681109246666078, 0.7116559215755541, 0.7116559215755541]  
['QDA', 'LinearSVC', 'KNN', 'GaussianNB']
```



In conclusion, the models used perform worse than the tree based and deep neural networks. Therefore, they won't be tested with the ids2018 dataset.