

5_ML_Default_2018

June 17, 2024

1 Machine Learning Models on the IDS 2018

In this notebook, decision tree and random forest based machine learning algorithms are applied to the ids2018 dataset. Several methods for resolving the class imbalance are tested. Decision tree algorithms were chosen for their effectiveness and the training time which were better than other machine learning models. RT and RF based algorithms performed better in the preliminary experiments

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import glob
import os
import xgboost as xgb
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, average_precision_score, \
    ↪make_scorer, precision_score, accuracy_score, confusion_matrix, \
    ↪recall_score, f1_score, roc_auc_score
from notebook_utils import load_sample_dataset_2018
%matplotlib inline
%load_ext autoreload
%autoreload 2

file_path = r"..\\CIC-IDS-2018\\Processed Traffic Data for ML Algorithms"

attack_labels = {
    0: 'Benign',
    1: 'Bot',
    2: 'Brute Force -Web',
    3: 'Brute Force -XSS',
    4: 'DDOS attack-HOIC',
    5: 'DDOS attack-LOIC-UDP',
    6: 'DDoS attacks-LOIC-HTTP',
    7: 'DoS attacks-GoldenEye',
```

```

8: 'DoS attacks-Hulk',
9: 'DoS attacks-SlowHTTPTest',
10: 'DoS attacks-Slowloris',
11: 'FTP-BruteForce',
12: 'Infiltration',
13: 'SQL Injection',
14: 'SSH-Bruteforce'
}

df = load_sample_dataset_2018(file_path)

```

Processed 1/10 files.
 Processed 2/10 files.
 Processed 3/10 files.
 Processed 4/10 files.
 Processed 5/10 files.
 Processed 6/10 files.
 Processed 7/10 files.
 Processed 8/10 files.
 Processed 9/10 files.
 Processed 10/10 files.

Creating is_attack column...

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1623303 entries, 0 to 1623302

Data columns (total 81 columns):

#	Column	Non-Null Count	Dtype
0	dst_port	1623295 non-null	float64
1	protocol	1623295 non-null	float64
2	timestamp	0 non-null	float64
3	flow_duration	1623295 non-null	float64
4	tot_fwd_pkts	1623295 non-null	float64
5	tot_bwd_pkts	1623295 non-null	float64
6	totlen_fwd_pkts	1623295 non-null	float64
7	totlen_bwd_pkts	1623295 non-null	float64
8	fwd_pkt_len_max	1623295 non-null	float64
9	fwd_pkt_len_min	1623295 non-null	float64
10	fwd_pkt_len_mean	1623295 non-null	float64
11	fwd_pkt_len_std	1623295 non-null	float64
12	bwd_pkt_len_max	1623295 non-null	float64
13	bwd_pkt_len_min	1623295 non-null	float64
14	bwd_pkt_len_mean	1623295 non-null	float64
15	bwd_pkt_len_std	1623295 non-null	float64
16	flow_byts_s	1617377 non-null	float64
17	flow_pkts_s	1623295 non-null	float64
18	flow_iat_mean	1623295 non-null	float64
19	flow_iat_std	1623295 non-null	float64
20	flow_iat_max	1623295 non-null	float64

21	flow_iat_min	1623295	non-null	float64
22	fwd_iat_tot	1623295	non-null	float64
23	fwd_iat_mean	1623295	non-null	float64
24	fwd_iat_std	1623295	non-null	float64
25	fwd_iat_max	1623295	non-null	float64
26	fwd_iat_min	1623295	non-null	float64
27	bwd_iat_tot	1623295	non-null	float64
28	bwd_iat_mean	1623295	non-null	float64
29	bwd_iat_std	1623295	non-null	float64
30	bwd_iat_max	1623295	non-null	float64
31	bwd_iat_min	1623295	non-null	float64
32	fwd_psh_flags	1623295	non-null	float64
33	bwd_psh_flags	1623295	non-null	float64
34	fwd_urg_flags	1623295	non-null	float64
35	bwd_urg_flags	1623295	non-null	float64
36	fwd_header_len	1623295	non-null	float64
37	bwd_header_len	1623295	non-null	float64
38	fwd_pkts_s	1623295	non-null	float64
39	bwd_pkts_s	1623295	non-null	float64
40	pkt_len_min	1623295	non-null	float64
41	pkt_len_max	1623295	non-null	float64
42	pkt_len_mean	1623295	non-null	float64
43	pkt_len_std	1623295	non-null	float64
44	pkt_len_var	1623295	non-null	float64
45	fin_flag_cnt	1623295	non-null	float64
46	syn_flag_cnt	1623295	non-null	float64
47	rst_flag_cnt	1623295	non-null	float64
48	psh_flag_cnt	1623295	non-null	float64
49	ack_flag_cnt	1623295	non-null	float64
50	urg_flag_cnt	1623295	non-null	float64
51	cwe_flag_count	1623295	non-null	float64
52	ece_flag_cnt	1623295	non-null	float64
53	down_up_ratio	1623295	non-null	float64
54	pkt_size_avg	1623295	non-null	float64
55	fwd_seg_size_avg	1623295	non-null	float64
56	bwd_seg_size_avg	1623295	non-null	float64
57	fwd_byts_b_avg	1623295	non-null	float64
58	fwd_pkts_b_avg	1623295	non-null	float64
59	fwd_blk_rate_avg	1623295	non-null	float64
60	bwd_byts_b_avg	1623295	non-null	float64
61	bwd_pkts_b_avg	1623295	non-null	float64
62	bwd_blk_rate_avg	1623295	non-null	float64
63	subflow_fwd_pkts	1623295	non-null	float64
64	subflow_fwd_byts	1623295	non-null	float64
65	subflow_bwd_pkts	1623295	non-null	float64
66	subflow_bwd_byts	1623295	non-null	float64
67	init_fwd_win_byts	1623295	non-null	float64
68	init_bwd_win_byts	1623295	non-null	float64

```

69 fwd_act_data_pkts 1623295 non-null float64
70 fwd_seg_size_min 1623295 non-null float64
71 active_mean      1623295 non-null float64
72 active_std       1623295 non-null float64
73 active_max       1623295 non-null float64
74 active_min       1623295 non-null float64
75 idle_mean        1623295 non-null float64
76 idle_std         1623295 non-null float64
77 idle_max         1623295 non-null float64
78 idle_min         1623295 non-null float64
79 label            1623303 non-null category
80 is_attack        1623303 non-null int64
dtypes: category(1), float64(79), int64(1)
memory usage: 992.3 MB
None

```

1.1 Preparing the Dataset

1.1.1 Check for invalid values

```

[2]: # Select only numeric columns
numeric_columns = df.select_dtypes(include=[np.number]).columns
# Identify columns with NaN, infinity, or negative values
nan_columns = df[numeric_columns].columns[df[numeric_columns].isna().any()]
inf_columns = df[numeric_columns].columns[np.isinf(df[numeric_columns]).any()]
neg_columns = df[numeric_columns].columns[(df[numeric_columns] < 0).any()]
print("Columns with NaN values:", nan_columns.tolist())
print("Columns with infinite values:", inf_columns.tolist())
print("Columns with negative values:", neg_columns.tolist())
# Calculate the percentage of NaN, infinite, and negative values
nan_percentage = df[nan_columns].isna().mean() * 100
# nan_percentage = nan_percentage[nan_percentage > 1]
inf_percentage = df[inf_columns].map(lambda x: np.isinf(x)).mean() * 100
neg_percentage = df[neg_columns].map(lambda x: x < 0).mean() * 100
print("Percentage of NaN values in each column:\n", nan_percentage)
print("Percentage of infinite values in each column:\n", inf_percentage)
print("Percentage of negative values in each column:\n", neg_percentage)

```

```

Columns with NaN values: ['dst_port', 'protocol', 'timestamp', 'flow_duration',
'tot_fwd_pkts', 'tot_bwd_pkts', 'totlen_fwd_pkts', 'totlen_bwd_pkts',
'fwd_pkt_len_max', 'fwd_pkt_len_min', 'fwd_pkt_len_mean', 'fwd_pkt_len_std',
'bwd_pkt_len_max', 'bwd_pkt_len_min', 'bwd_pkt_len_mean', 'bwd_pkt_len_std',
'flow_byts_s', 'flow_pkts_s', 'flow_iat_mean', 'flow_iat_std', 'flow_iat_max',
'flow_iat_min', 'fwd_iat_tot', 'fwd_iat_mean', 'fwd_iat_std', 'fwd_iat_max',
'fwd_iat_min', 'bwd_iat_tot', 'bwd_iat_mean', 'bwd_iat_std', 'bwd_iat_max',
'bwd_iat_min', 'fwd_psh_flags', 'bwd_psh_flags', 'fwd_urg_flags',
'bwd_urg_flags', 'fwd_header_len', 'bwd_header_len', 'fwd_pkts_s', 'bwd_pkts_s',
'pkt_len_min', 'pkt_len_max', 'pkt_len_mean', 'pkt_len_std', 'pkt_len_var',
'fin_flag_cnt', 'syn_flag_cnt', 'rst_flag_cnt', 'psh_flag_cnt', 'ack_flag_cnt',

```

```

'urg_flag_cnt', 'cwe_flag_count', 'ece_flag_cnt', 'down_up_ratio',
'pkt_size_avg', 'fwd_seg_size_avg', 'bwd_seg_size_avg', 'fwd_byts_b_avg',
'fwd_pkts_b_avg', 'fwd_blk_rate_avg', 'bwd_byts_b_avg', 'bwd_pkts_b_avg',
'bwd_blk_rate_avg', 'subflow_fwd_pkts', 'subflow_fwd_byts', 'subflow_bwd_pkts',
'subflow_bwd_byts', 'init_fwd_win_byts', 'init_bwd_win_byts',
'fwd_act_data_pkts', 'fwd_seg_size_min', 'active_mean', 'active_std',
'active_max', 'active_min', 'idle_mean', 'idle_std', 'idle_max', 'idle_min']
Columns with infinite values: ['flow_byts_s', 'flow_pkts_s']
Columns with negative values: ['flow_duration', 'flow_pkts_s', 'flow_iat_mean',
'flow_iat_min', 'fwd_iat_tot', 'fwd_iat_mean', 'fwd_iat_min',
'init_fwd_win_byts', 'init_bwd_win_byts']
Percentage of NaN values in each column:
  dst_port      0.000493
  protocol      0.000493
  timestamp     100.000000
  flow_duration  0.000493
  tot_fwd_pkts  0.000493
  ...
  active_min    0.000493
  idle_mean     0.000493
  idle_std      0.000493
  idle_max      0.000493
  idle_min      0.000493
Length: 79, dtype: float64
Percentage of infinite values in each column:
  flow_byts_s    0.219182
  flow_pkts_s    0.583746
dtype: float64
Percentage of negative values in each column:
  flow_duration    0.000123
  flow_pkts_s      0.000123
  flow_iat_mean    0.000123
  flow_iat_min     0.000123
  fwd_iat_tot      0.000123
  fwd_iat_mean     0.000123
  fwd_iat_min      0.000123
  init_fwd_win_byts 27.295235
  init_bwd_win_byts 50.867888
dtype: float64

```

For negative values, 2 columns have an extremely high percentage of negative values. We choose to drop the features “init_win_bytes_forward” and “init_win_bytes_backward” as the source of the negative sign is unknown. For the rest of relevant features, the percentages of negative, infinite or are low so the rows are dropped.

```

[3]: def replace_invalid(df):
      # Select only numeric columns
      numeric_columns = df.select_dtypes(include=[np.number]).columns

```

```

# Identify columns with NaN, infinite, or negative values
nan_columns = df[numeric_columns].columns[df[numeric_columns].isna().any()]
inf_columns = df[numeric_columns].columns[np.isinf(df[numeric_columns]).
↳any()]
neg_columns = df[numeric_columns].columns[(df[numeric_columns] < 0).any()]
# Drop rows with NaN values (low percentage of NaN values)
# df = df.dropna(subset=nan_columns)
# Drop rows with infinite values (assuming low percentage)
for col in inf_columns:
    df = df[np.isfinite(df[col])]
# Drop columns with a high percentage of negative values
columns_to_drop = ['init_fwd_win_byts', 'init_bwd_win_byts']
df = df.drop(columns=columns_to_drop)
# Drop rows with negative values in the remaining columns
remaining_neg_columns = [col for col in neg_columns if col not in_
↳columns_to_drop]
for col in remaining_neg_columns:
    df = df[df[col] >= 0]
return df

```

```
[4]: df = replace_invalid(df)
```

```

[5]: X = df.iloc[:, 0:76]
Y = df[["label", "is_attack", "label_code"]]

X.info()
Y.info()
print(Y.label.value_counts())

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 1613823 entries, 0 to 1250804
```

```
Data columns (total 76 columns):
```

#	Column	Non-Null Count	Dtype
0	dst_port	1613823 non-null	float64
1	protocol	1613823 non-null	float64
2	timestamp	0 non-null	float64
3	flow_duration	1613823 non-null	float64
4	tot_fwd_pkts	1613823 non-null	float64
5	tot_bwd_pkts	1613823 non-null	float64
6	totlen_fwd_pkts	1613823 non-null	float64
7	totlen_bwd_pkts	1613823 non-null	float64
8	fwd_pkt_len_max	1613823 non-null	float64
9	fwd_pkt_len_min	1613823 non-null	float64
10	fwd_pkt_len_mean	1613823 non-null	float64
11	fwd_pkt_len_std	1613823 non-null	float64
12	bwd_pkt_len_max	1613823 non-null	float64
13	bwd_pkt_len_min	1613823 non-null	float64

14	bwd_pkt_len_mean	1613823	non-null	float64
15	bwd_pkt_len_std	1613823	non-null	float64
16	flow_byts_s	1613823	non-null	float64
17	flow_pkts_s	1613823	non-null	float64
18	flow_iat_mean	1613823	non-null	float64
19	flow_iat_std	1613823	non-null	float64
20	flow_iat_max	1613823	non-null	float64
21	flow_iat_min	1613823	non-null	float64
22	fwd_iat_tot	1613823	non-null	float64
23	fwd_iat_mean	1613823	non-null	float64
24	fwd_iat_std	1613823	non-null	float64
25	fwd_iat_max	1613823	non-null	float64
26	fwd_iat_min	1613823	non-null	float64
27	bwd_iat_tot	1613823	non-null	float64
28	bwd_iat_mean	1613823	non-null	float64
29	bwd_iat_std	1613823	non-null	float64
30	bwd_iat_max	1613823	non-null	float64
31	bwd_iat_min	1613823	non-null	float64
32	fwd_psh_flags	1613823	non-null	float64
33	bwd_psh_flags	1613823	non-null	float64
34	fwd_urg_flags	1613823	non-null	float64
35	bwd_urg_flags	1613823	non-null	float64
36	fwd_header_len	1613823	non-null	float64
37	bwd_header_len	1613823	non-null	float64
38	fwd_pkts_s	1613823	non-null	float64
39	bwd_pkts_s	1613823	non-null	float64
40	pkt_len_min	1613823	non-null	float64
41	pkt_len_max	1613823	non-null	float64
42	pkt_len_mean	1613823	non-null	float64
43	pkt_len_std	1613823	non-null	float64
44	pkt_len_var	1613823	non-null	float64
45	fin_flag_cnt	1613823	non-null	float64
46	syn_flag_cnt	1613823	non-null	float64
47	rst_flag_cnt	1613823	non-null	float64
48	psh_flag_cnt	1613823	non-null	float64
49	ack_flag_cnt	1613823	non-null	float64
50	urg_flag_cnt	1613823	non-null	float64
51	cwe_flag_count	1613823	non-null	float64
52	ece_flag_cnt	1613823	non-null	float64
53	down_up_ratio	1613823	non-null	float64
54	pkt_size_avg	1613823	non-null	float64
55	fwd_seg_size_avg	1613823	non-null	float64
56	bwd_seg_size_avg	1613823	non-null	float64
57	fwd_byts_b_avg	1613823	non-null	float64
58	fwd_pkts_b_avg	1613823	non-null	float64
59	fwd_blk_rate_avg	1613823	non-null	float64
60	bwd_byts_b_avg	1613823	non-null	float64
61	bwd_pkts_b_avg	1613823	non-null	float64

```

62 bwd_blk_rate_avg      1613823 non-null float64
63 subflow_fwd_pkts      1613823 non-null float64
64 subflow_fwd_byts      1613823 non-null float64
65 subflow_bwd_pkts      1613823 non-null float64
66 subflow_bwd_byts      1613823 non-null float64
67 fwd_act_data_pkts      1613823 non-null float64
68 fwd_seg_size_min       1613823 non-null float64
69 active_mean            1613823 non-null float64
70 active_std             1613823 non-null float64
71 active_max            1613823 non-null float64
72 active_min            1613823 non-null float64
73 idle_mean             1613823 non-null float64
74 idle_std              1613823 non-null float64
75 idle_max              1613823 non-null float64

```

dtypes: float64(76)

memory usage: 948.1 MB

<class 'pandas.core.frame.DataFrame'>

Index: 1613823 entries, 0 to 1250804

Data columns (total 3 columns):

#	Column	Non-Null Count	Dtype
0	label	1613823 non-null	category
1	is_attack	1613823 non-null	int64
2	label_code	1613823 non-null	int32

dtypes: category(1), int32(1), int64(1)

memory usage: 32.3 MB

label

Benign	1338596
DDoS attack-HOIC	68817
DDoS attacks-LOIC-HTTP	57678
DoS attacks-Hulk	46307
Bot	28759
FTP-BruteForce	19396
SSH-Bruteforce	18759
Infiltration	15952
DoS attacks-SlowHTTPTest	14029
DoS attacks-GoldenEye	4215
DoS attacks-Slowloris	1055
DDoS attack-LOIC-UDP	168
Brute Force -Web	53
Brute Force -XSS	27
SQL Injection	12
Label	0

Name: count, dtype: int64

1.2 Feature Selection

First, the columns with no variance are dropped as they have no impact on the target variables.


```
[6]: stats = X.describe()
std = stats.loc["std"]
features_no_var = std[std == 0.0].index
# Exclude non-numeric columns (e.g., categorical columns) from the features_
↳with zero variance
features_no_var_numeric = [col for col in features_no_var if col in X.
↳select_dtypes(include=[np.number]).columns]
print(features_no_var_numeric)
```

```
['bwd_psh_flags', 'bwd_urg_flags', 'fwd_byts_b_avg', 'fwd_pkts_b_avg',
'fwd_blk_rate_avg', 'bwd_byts_b_avg', 'bwd_pkts_b_avg', 'bwd_blk_rate_avg']
```

```
[7]: X = X.drop(columns=features_no_var)
X = X.drop(columns=['dst_port', 'timestamp'])
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1613823 entries, 0 to 1250804
Data columns (total 66 columns):
#   Column                Non-Null Count  Dtype
---  -
0   protocol              1613823 non-null float64
1   flow_duration         1613823 non-null float64
2   tot_fwd_pkts          1613823 non-null float64
3   tot_bwd_pkts          1613823 non-null float64
4   totlen_fwd_pkts       1613823 non-null float64
5   totlen_bwd_pkts       1613823 non-null float64
6   fwd_pkt_len_max       1613823 non-null float64
7   fwd_pkt_len_min       1613823 non-null float64
8   fwd_pkt_len_mean      1613823 non-null float64
9   fwd_pkt_len_std       1613823 non-null float64
10  bwd_pkt_len_max        1613823 non-null float64
11  bwd_pkt_len_min        1613823 non-null float64
12  bwd_pkt_len_mean       1613823 non-null float64
13  bwd_pkt_len_std        1613823 non-null float64
14  flow_byts_s            1613823 non-null float64
15  flow_pkts_s            1613823 non-null float64
16  flow_iat_mean          1613823 non-null float64
17  flow_iat_std           1613823 non-null float64
18  flow_iat_max           1613823 non-null float64
19  flow_iat_min           1613823 non-null float64
20  fwd_iat_tot            1613823 non-null float64
21  fwd_iat_mean           1613823 non-null float64
22  fwd_iat_std            1613823 non-null float64
23  fwd_iat_max            1613823 non-null float64
24  fwd_iat_min            1613823 non-null float64
25  bwd_iat_tot            1613823 non-null float64
26  bwd_iat_mean           1613823 non-null float64
```

27	bwd_iat_std	1613823	non-null	float64
28	bwd_iat_max	1613823	non-null	float64
29	bwd_iat_min	1613823	non-null	float64
30	fwd_psh_flags	1613823	non-null	float64
31	fwd_urg_flags	1613823	non-null	float64
32	fwd_header_len	1613823	non-null	float64
33	bwd_header_len	1613823	non-null	float64
34	fwd_pkts_s	1613823	non-null	float64
35	bwd_pkts_s	1613823	non-null	float64
36	pkt_len_min	1613823	non-null	float64
37	pkt_len_max	1613823	non-null	float64
38	pkt_len_mean	1613823	non-null	float64
39	pkt_len_std	1613823	non-null	float64
40	pkt_len_var	1613823	non-null	float64
41	fin_flag_cnt	1613823	non-null	float64
42	syn_flag_cnt	1613823	non-null	float64
43	rst_flag_cnt	1613823	non-null	float64
44	psh_flag_cnt	1613823	non-null	float64
45	ack_flag_cnt	1613823	non-null	float64
46	urg_flag_cnt	1613823	non-null	float64
47	cwe_flag_count	1613823	non-null	float64
48	ece_flag_cnt	1613823	non-null	float64
49	down_up_ratio	1613823	non-null	float64
50	pkt_size_avg	1613823	non-null	float64
51	fwd_seg_size_avg	1613823	non-null	float64
52	bwd_seg_size_avg	1613823	non-null	float64
53	subflow_fwd_pkts	1613823	non-null	float64
54	subflow_fwd_byts	1613823	non-null	float64
55	subflow_bwd_pkts	1613823	non-null	float64
56	subflow_bwd_byts	1613823	non-null	float64
57	fwd_act_data_pkts	1613823	non-null	float64
58	fwd_seg_size_min	1613823	non-null	float64
59	active_mean	1613823	non-null	float64
60	active_std	1613823	non-null	float64
61	active_max	1613823	non-null	float64
62	active_min	1613823	non-null	float64
63	idle_mean	1613823	non-null	float64
64	idle_std	1613823	non-null	float64
65	idle_max	1613823	non-null	float64

dtypes: float64(66)

memory usage: 824.9 MB

1.2.1 Remove collinear variables

```
[8]: def correlation_feature_selection(df, threshold=0.95):
      corr_matrix = df.corr().abs()
```

```

    upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).
↳astype(bool))
    to_drop = [column for column in upper.columns if any(upper[column] >↳
↳threshold)]
    return df.drop(columns=to_drop)
X = correlation_feature_selection(X)
X.info()

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 1613823 entries, 0 to 1250804
```

```
Data columns (total 41 columns):
```

#	Column	Non-Null Count	Dtype
0	protocol	1613823 non-null	float64
1	flow_duration	1613823 non-null	float64
2	tot_fwd_pkts	1613823 non-null	float64
3	tot_bwd_pkts	1613823 non-null	float64
4	fwd_pkt_len_max	1613823 non-null	float64
5	fwd_pkt_len_min	1613823 non-null	float64
6	fwd_pkt_len_mean	1613823 non-null	float64
7	bwd_pkt_len_max	1613823 non-null	float64
8	bwd_pkt_len_min	1613823 non-null	float64
9	bwd_pkt_len_mean	1613823 non-null	float64
10	flow_byts_s	1613823 non-null	float64
11	flow_pkts_s	1613823 non-null	float64
12	flow_iat_mean	1613823 non-null	float64
13	flow_iat_std	1613823 non-null	float64
14	flow_iat_max	1613823 non-null	float64
15	fwd_iat_std	1613823 non-null	float64
16	bwd_iat_tot	1613823 non-null	float64
17	bwd_iat_mean	1613823 non-null	float64
18	bwd_iat_std	1613823 non-null	float64
19	bwd_iat_max	1613823 non-null	float64
20	bwd_iat_min	1613823 non-null	float64
21	fwd_psh_flags	1613823 non-null	float64
22	fwd_urg_flags	1613823 non-null	float64
23	fwd_pkts_s	1613823 non-null	float64
24	bwd_pkts_s	1613823 non-null	float64
25	pkt_len_min	1613823 non-null	float64
26	pkt_len_mean	1613823 non-null	float64
27	pkt_len_var	1613823 non-null	float64
28	fin_flag_cnt	1613823 non-null	float64
29	rst_flag_cnt	1613823 non-null	float64
30	psh_flag_cnt	1613823 non-null	float64
31	ack_flag_cnt	1613823 non-null	float64
32	urg_flag_cnt	1613823 non-null	float64
33	down_up_ratio	1613823 non-null	float64

```

34 fwd_seg_size_min 1613823 non-null float64
35 active_mean      1613823 non-null float64
36 active_std       1613823 non-null float64
37 active_max       1613823 non-null float64
38 active_min       1613823 non-null float64
39 idle_mean        1613823 non-null float64
40 idle_std         1613823 non-null float64
dtypes: float64(41)
memory usage: 517.1 MB

```

1.2.2 Information Gain Selection

```

[9]: from sklearn.feature_selection import mutual_info_classif
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import RandomOverSampler

def oversample_minority_classes(X, Y, sample_size=1000):
    y=Y["label_code"]
    ros = RandomOverSampler(random_state=42)
    X_resampled, y_resampled = ros.fit_resample(X, y)
    # Create a subset of the oversampled data
    X_sample, _, y_sample, _ = train_test_split(X_resampled, y_resampled,
    ↪train_size=sample_size, stratify=y_resampled, random_state=42)
    return X_sample, y_sample

def information_gain_feature_selection(X, Y, sample_size=1000):
    # Create an oversampled subset of the data
    X_sample, y_sample = oversample_minority_classes(X, Y, sample_size)
    # Create is_attack column based on label_code
    y_sample = (y_sample != 0).astype(int)
    # Perform feature selection on the oversampled subset
    info_gain = mutual_info_classif(X_sample, y_sample)
    info_gain_df = pd.DataFrame({'Feature': X.columns, 'Information Gain':
    ↪info_gain})
    info_gain_df = info_gain_df.sort_values(by='Information Gain',
    ↪ascending=False)
    print(info_gain_df)
    selected_features = info_gain_df[info_gain_df['Information Gain'] >
    ↪0]['Feature'].tolist()
    return selected_features

# Determine the selected features using the oversampled subset
selected_features = information_gain_feature_selection(X, Y)

# Apply the selected features to the main dataset
X = X[selected_features]

```

```
# Display information about the selected features
X.info()
```

	Feature	Information Gain
7	bwd_pkt_len_max	0.098128
27	pkt_len_var	0.084071
4	fwd_pkt_len_max	0.079047
6	fwd_pkt_len_mean	0.075914
9	bwd_pkt_len_mean	0.064383
2	tot_fwd_pkts	0.051284
26	pkt_len_mean	0.050098
1	flow_duration	0.046261
24	bwd_pkts_s	0.045107
14	flow_iat_max	0.043121
13	flow_iat_std	0.040591
34	fwd_seg_size_min	0.040315
8	bwd_pkt_len_min	0.038632
10	flow_byts_s	0.035294
18	bwd_iat_std	0.032327
20	bwd_iat_min	0.030660
23	fwd_pkts_s	0.029708
5	fwd_pkt_len_min	0.027187
12	flow_iat_mean	0.026420
11	flow_pkts_s	0.025621
15	fwd_iat_std	0.023797
3	tot_bwd_pkts	0.022471
16	bwd_iat_tot	0.021173
25	pkt_len_min	0.019351
19	bwd_iat_max	0.015935
17	bwd_iat_mean	0.015099
38	active_min	0.012552
0	protocol	0.011281
37	active_max	0.005979
36	active_std	0.004185
32	urg_flag_cnt	0.001447
39	idle_mean	0.000311
35	active_mean	0.000158
28	fin_flag_cnt	0.000000
29	rst_flag_cnt	0.000000
30	psh_flag_cnt	0.000000
31	ack_flag_cnt	0.000000
33	down_up_ratio	0.000000
22	fwd_urg_flags	0.000000
21	fwd_psh_flags	0.000000
40	idle_std	0.000000

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 1613823 entries, 0 to 1250804
```

```
Data columns (total 33 columns):
```

#	Column	Non-Null Count	Dtype
0	bwd_pkt_len_max	1613823 non-null	float64
1	pkt_len_var	1613823 non-null	float64
2	fwd_pkt_len_max	1613823 non-null	float64
3	fwd_pkt_len_mean	1613823 non-null	float64
4	bwd_pkt_len_mean	1613823 non-null	float64
5	tot_fwd_pkts	1613823 non-null	float64
6	pkt_len_mean	1613823 non-null	float64
7	flow_duration	1613823 non-null	float64
8	bwd_pkts_s	1613823 non-null	float64
9	flow_iat_max	1613823 non-null	float64
10	flow_iat_std	1613823 non-null	float64
11	fwd_seg_size_min	1613823 non-null	float64
12	bwd_pkt_len_min	1613823 non-null	float64
13	flow_byts_s	1613823 non-null	float64
14	bwd_iat_std	1613823 non-null	float64
15	bwd_iat_min	1613823 non-null	float64
16	fwd_pkts_s	1613823 non-null	float64
17	fwd_pkt_len_min	1613823 non-null	float64
18	flow_iat_mean	1613823 non-null	float64
19	flow_pkts_s	1613823 non-null	float64
20	fwd_iat_std	1613823 non-null	float64
21	tot_bwd_pkts	1613823 non-null	float64
22	bwd_iat_tot	1613823 non-null	float64
23	pkt_len_min	1613823 non-null	float64
24	bwd_iat_max	1613823 non-null	float64
25	bwd_iat_mean	1613823 non-null	float64
26	active_min	1613823 non-null	float64
27	protocol	1613823 non-null	float64
28	active_max	1613823 non-null	float64
29	active_std	1613823 non-null	float64
30	urg_flag_cnt	1613823 non-null	float64
31	idle_mean	1613823 non-null	float64
32	active_mean	1613823 non-null	float64

dtypes: float64(33)

memory usage: 418.6 MB

1.3 Split Dataset

The dataset is split into a training set and a testing set with a ratio of 0.8/0.2. The dataset is stratified according to the label to have an equal representation of all classes in the 2 subsets.

```
[10]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
↳ stratify=Y.label)
```

```
[11]: Y_train.label.value_counts()
```

```
[11]: label
      Benign                1070876
      DDOS attack-HOIC      55054
      DDOS attacks-LOIC-HTTP 46142
      DoS attacks-Hulk      37046
      Bot                   23007
      FTP-BruteForce        15517
      SSH-Bruteforce        15007
      Infiltration          12762
      DoS attacks-SlowHTTPTest 11223
      DoS attacks-GoldenEye  3372
      DoS attacks-Slowloris  844
      DDOS attack-LOIC-UDP   134
      Brute Force -Web       42
      Brute Force -XSS       22
      SQL Injection          10
      Label                  0
      Name: count, dtype: int64
```

```
[12]: Y_test.label.value_counts()
```

```
[12]: label
      Benign                267720
      DDOS attack-HOIC      13763
      DDOS attacks-LOIC-HTTP 11536
      DoS attacks-Hulk      9261
      Bot                   5752
      FTP-BruteForce        3879
      SSH-Bruteforce        3752
      Infiltration          3190
      DoS attacks-SlowHTTPTest 2806
      DoS attacks-GoldenEye  843
      DoS attacks-Slowloris  211
      DDOS attack-LOIC-UDP   34
      Brute Force -Web       11
      Brute Force -XSS       5
      SQL Injection          2
      Label                  0
      Name: count, dtype: int64
```

```
[13]: benign_percentage = len(Y_train.label[Y_train["label"]=="BENIGN"])/len(Y_train)
      print('Percentage of benign samples: %.4f' % benign_percentage)
      print(Y_train.is_attack.value_counts())
```

```
Percentage of benign samples: 0.0000
is_attack
0    1070876
1    220182
```

Name: count, dtype: int64

1.4 Machine Learning Classifiers with Default Hyperparameters

1.4.1 Helper Functions

```
[14]: import joblib

def save_model(model, model_name):
    file_path = f'models/{model_name}.pkl'
    joblib.dump(model, file_path)
    print(f'Model saved to {file_path}')

def load_model(model_name):
    file_path = f'models/{model_name}.pkl'
    model = joblib.load(file_path)
    print(f'Model loaded from {file_path}')
    return model

os.makedirs('models', exist_ok=True)

[15]: def plot_confusion_matrix(model_name, Y_true, Y_pred, labels=["Benign",
    ↪ "Attack"]):
    matrix = confusion_matrix(Y_true.is_attack, Y_pred)
    plt.figure(figsize=(8, 6))
    sns.heatmap(matrix, annot=True, cmap='Blues', fmt='d', xticklabels=labels,
    ↪ yticklabels=labels)
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.title(f'Confusion Matrix for {model_name}')
    plt.show()

def metrics_report(dataset_type, y_true, y_predict, print_avg=True):
    print(f"Classification Report ({dataset_type}):")
    print(classification_report(y_true, y_predict, digits=4))
    accuracy = accuracy_score(y_true, y_predict)
    precision = precision_score(y_true, y_predict, average='weighted')
    recall = recall_score(y_true, y_predict, average='weighted')
    f1 = f1_score(y_true, y_predict, average='weighted')
    auc = roc_auc_score(y_true, y_predict)
    print("Accuracy:", accuracy)
    print("Precision:", precision)
    print("Recall:", recall)
    print("F1 Score:", f1)
    print("AUC:", auc)
    return {"accuracy": accuracy, "precision": precision, "recall": recall,
    ↪ "f1": f1, "auc": auc}
```



```

def calculate_metrics_by_label(y_true, y_pred, labels):
    results = []
    unique_labels = labels.unique()
    for label in unique_labels:
        indices = labels == label
        accuracy = accuracy_score(y_true[indices], y_pred[indices])
        results.append({
            'Label': label,
            'Accuracy': accuracy,
        })
    return pd.DataFrame(results)

```

```

[16]: def test_metrics(model_name, model, dataset_type, scaler):
    # Predict and evaluate on the test set
    print(f"{model_name} with {dataset_type} Test Set Performance")
    Y_pred = model.predict(scaler.transform(X_test))
    metrics = metrics_report(f"Test {model_name} ({dataset_type})", Y_test.
↪is_attack, Y_pred)
    plot_confusion_matrix(f"{model_name} ({dataset_type})", Y_test, Y_pred)
    # Calculate metrics by label
    metrics_by_label = calculate_metrics_by_label(Y_test.is_attack, Y_pred,
↪Y_test.label)
    metrics_by_label['Method'] = dataset_type
    print(f"Metrics by Label ({dataset_type}):")
    print(metrics_by_label)
    return metrics, metrics_by_label

```

```

[17]: def plot_overall_accuracy(metrics):
    methods = ['original', 'random', 'smote', 'adasyn']
    overall_accuracies = []

    # Extract overall accuracy for each method
    for method in methods:
        overall_accuracies.append(metrics[method][0]['accuracy'])

    # Plotting the overall accuracies
    plt.figure(figsize=(10, 6))
    bars = plt.bar(methods, overall_accuracies, color=['blue', 'orange',
↪'green', 'red'])
    plt.title('Overall Accuracy by Method')
    plt.xlabel('Method')
    plt.ylabel('Accuracy')
    plt.ylim(0.9, 1)
    plt.grid(True)

    # Display the values on each bar
    for bar in bars:

```

```

        yval = bar.get_height()
        plt.text(bar.get_x() + bar.get_width() / 2.0, yval, f'{yval:.3f}',
        ↪ha='center', va='bottom')

plt.show()

```

1.4.2 Resampling methods

```

[18]: from imblearn.over_sampling import RandomOverSampler, SMOTE, ADASYN

def resample_dataset(X, Y, min_samples, attack_labels, technique='smote'):
    Y = Y.drop(columns=['label'])
    combined = pd.concat([X, Y], axis=1)
    counts = Y['label_code'].value_counts()
    samples_number = {i: max(counts[i], min_samples) for i in np.
    ↪unique(Y['label_code'])}
    combined_array = combined.values
    y_array = Y['label_code'].values

    if technique == 'random':
        resampler = RandomOverSampler(random_state=42,
    ↪sampling_strategy=samples_number)
    elif technique == 'smote':
        resampler = SMOTE(random_state=42, sampling_strategy=samples_number,
    ↪k_neighbors=5)
    elif technique == 'adasyn':
        resampler = ADASYN(random_state=42, sampling_strategy=samples_number)
    else:
        raise ValueError("Invalid resampling technique. Choose 'random',
    ↪'smote', or 'adasyn'.")

    resampled_array, y_resampled = resampler.fit_resample(combined_array,
    ↪y_array)
    X_resampled = resampled_array[:, :-Y.shape[1]]
    Y_resampled = resampled_array[:, -Y.shape[1]:]
    X_resampled_df = pd.DataFrame(X_resampled, columns=X.columns)
    Y_resampled_df = pd.DataFrame(Y_resampled, columns=Y.columns)
    Y_resampled_df['label'] = Y_resampled_df['label_code'].map(attack_labels)
    Y_resampled_df['label'] = Y_resampled_df['label'].astype('category')
    return X_resampled_df, Y_resampled_df

[19]: X_random_train, Y_random_train = resample_dataset(X_train, Y_train, 100000,
    ↪attack_labels, "random")
X_smote_train, Y_smote_train = resample_dataset(X_train, Y_train, 100000,
    ↪attack_labels, "smote")

```

```
X_adasyn_train, Y_adasyn_train = resample_dataset(X_train, Y_train, 100000,
↪attack_labels, "adasyn")
```

```
[20]: Y_train.label.value_counts()
```

```
[20]: label
      Benign                1070876
      DDOS attack-HOIC        55054
      DDoS attacks-LOIC-HTTP  46142
      DoS attacks-Hulk        37046
      Bot                    23007
      FTP-BruteForce          15517
      SSH-Bruteforce          15007
      Infiltration            12762
      DoS attacks-SlowHTTPTest 11223
      DoS attacks-GoldenEye    3372
      DoS attacks-Slowloris     844
      DDOS attack-LOIC-UDP     134
      Brute Force -Web         42
      Brute Force -XSS         22
      SQL Injection            10
      Label                    0
      Name: count, dtype: int64
```

```
[21]: Y_random_train.label.value_counts()
```

```
[21]: label
      Benign                1070876
      Bot                    100000
      Brute Force -Web       100000
      Brute Force -XSS       100000
      DDOS attack-HOIC       100000
      DDOS attack-LOIC-UDP   100000
      DDoS attacks-LOIC-HTTP 100000
      DoS attacks-GoldenEye   100000
      DoS attacks-Hulk        100000
      DoS attacks-SlowHTTPTest 100000
      DoS attacks-Slowloris   100000
      FTP-BruteForce          100000
      Infiltration            100000
      SSH-Bruteforce          100000
      Name: count, dtype: int64
```

```
[22]: Y_smote_train.label.value_counts()
```

```
[22]: label
      Benign                1070876
```

Bot	100000
Brute Force -Web	100000
Brute Force -XSS	100000
DDOS attack-HOIC	100000
DDOS attack-LOIC-UDP	100000
DDoS attacks-LOIC-HTTP	100000
DoS attacks-GoldenEye	100000
DoS attacks-Hulk	100000
DoS attacks-SlowHTTPTest	100000
DoS attacks-Slowloris	100000
FTP-BruteForce	100000
Infiltration	100000
SSH-Bruteforce	100000

Name: count, dtype: int64

```
[23]: Y_adasyn_train.label.value_counts()
```

```
[23]: label
```

Benign	1070876
Infiltration	102574
DoS attacks-GoldenEye	100245
DoS attacks-Slowloris	100008
Brute Force -XSS	100006
DDoS attacks-LOIC-HTTP	100003
SSH-Bruteforce	100000
FTP-BruteForce	99999
DDOS attack-LOIC-UDP	99998
DoS attacks-SlowHTTPTest	99997
Brute Force -Web	99994
Bot	99985
DDOS attack-HOIC	99550
DoS attacks-Hulk	98950

Name: count, dtype: int64

1.4.3 Scaling with the Standard Scaler

```
[24]: # Original X_train
scaler = StandardScaler()
scaler.fit(X_train)
# Random Oversampling
scaler_random = StandardScaler()
scaler_random.fit(X_random_train)
# SMOTE
scaler_smote = StandardScaler()
scaler_smote.fit(X_smote_train)
# ADASYN
scaler_adasyn = StandardScaler()
```

```
scaler_adasyn.fit(X_adasyn_train)
```

```
[24]: StandardScaler()
```

1.4.4 Decision Tree

```
[25]: dt_metrics = {}
```

```
[26]: decision_tree_model = DecisionTreeClassifier()  
decision_tree_model.fit(scaler.transform(X_train), Y_train.is_attack)
```

```
[26]: DecisionTreeClassifier()
```

```
[27]: # Predict and evaluate on the test set  
dt_metrics["original"] = test_metrics("Decision Tree", decision_tree_model,   
↪ "Original", scaler)
```

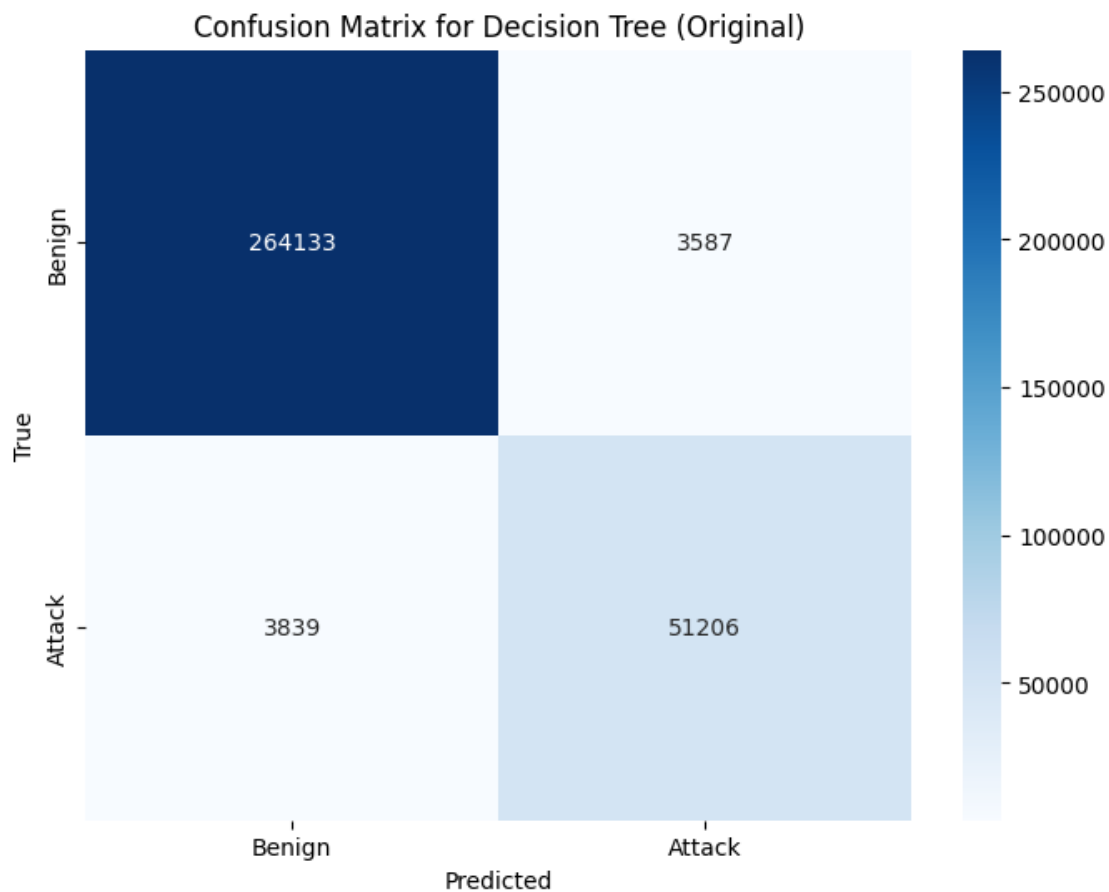
Decision Tree with Original Test Set Performance
Classification Report (Test Decision Tree (Original)):

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.9857	0.9866	0.9861	267720
1	0.9345	0.9303	0.9324	55045

accuracy			0.9770	322765
macro avg	0.9601	0.9584	0.9593	322765
weighted avg	0.9770	0.9770	0.9770	322765

Accuracy: 0.9769925487583846
Precision: 0.9769526222247139
Recall: 0.9769925487583846
F1 Score: 0.976971567476951
AUC: 0.9584293678967986



Metrics by Label (Original):

	Label	Accuracy	Method
0	Benign	0.986602	Original
1	DDoS attack-HOIC	0.951609	Original
2	DDoS attacks-LOIC-HTTP	0.981016	Original
3	DoS attacks-Hulk	0.995681	Original
4	DoS attacks-SlowHTTPTest	1.000000	Original
5	SSH-Bruteforce	0.999467	Original
6	FTP-BruteForce	1.000000	Original
7	Infiltration	0.101881	Original
8	DoS attacks-GoldenEye	0.998814	Original
9	Bot	0.994958	Original
10	DDoS attack-LOIC-UDP	1.000000	Original
11	DoS attacks-Slowloris	0.962085	Original
12	Brute Force -Web	0.545455	Original
13	Brute Force -XSS	0.400000	Original
14	SQL Injection	0.500000	Original

```
[28]: decision_tree_model_random = DecisionTreeClassifier()
      decision_tree_model_random.fit(scaler_random.transform(X_random_train),
      ↪Y_random_train.is_attack)
```


```
[28]: DecisionTreeClassifier()
```

```
[29]: # Predict and evaluate on the test set
      dt_metrics["random"] = test_metrics("Decision Tree",
      ↪decision_tree_model_random, "Random Oversampling", scaler_random)
```

Decision Tree with Random Oversampling Test Set Performance
 Classification Report (Test Decision Tree (Random Oversampling)):

	precision	recall	f1-score	support
0	0.9881	0.9724	0.9802	267720
1	0.8755	0.9429	0.9079	55045
accuracy			0.9674	322765
macro avg	0.9318	0.9577	0.9441	322765
weighted avg	0.9689	0.9674	0.9679	322765

Accuracy: 0.9673911359657955
 Precision: 0.9688707442936672
 Recall: 0.9673911359657955
 F1 Score: 0.9678657616700608
 AUC: 0.9576710549850029



Metrics by Label (Random Oversampling):

	Label	Accuracy	Method
0	Benign	0.972423	Random Oversampling
1	DDoS attack-HOIC	0.967812	Random Oversampling
2	DDoS attacks-LOIC-HTTP	0.981276	Random Oversampling
3	DoS attacks-Hulk	0.997408	Random Oversampling
4	DoS attacks-SlowHTTPTest	1.000000	Random Oversampling
5	SSH-Bruteforce	0.999733	Random Oversampling
6	FTP-BruteForce	1.000000	Random Oversampling
7	Infiltration	0.239498	Random Oversampling
8	DoS attacks-GoldenEye	1.000000	Random Oversampling
9	Bot	0.997392	Random Oversampling
10	DDoS attack-LOIC-UDP	1.000000	Random Oversampling
11	DoS attacks-Slowloris	0.952607	Random Oversampling
12	Brute Force -Web	0.545455	Random Oversampling
13	Brute Force -XSS	0.600000	Random Oversampling
14	SQL Injection	1.000000	Random Oversampling


```
[30]: decision_tree_model_smote = DecisionTreeClassifier()
      decision_tree_model_smote.fit(scaler_smote.transform(X_smote_train),
      ↪Y_smote_train.is_attack)
```

```
[30]: DecisionTreeClassifier()
```

```
[31]: # Predict and evaluate on the test set
      dt_metrics["smote"] = test_metrics("Decision Tree", decision_tree_model_smote,
      ↪"SMOTE", scaler_smote)
```

Decision Tree with SMOTE Test Set Performance

Classification Report (Test Decision Tree (SMOTE)):

	precision	recall	f1-score	support
0	0.9881	0.9764	0.9822	267720
1	0.8914	0.9428	0.9164	55045
accuracy			0.9707	322765
macro avg	0.9398	0.9596	0.9493	322765
weighted avg	0.9716	0.9707	0.9710	322765


Accuracy: 0.9706566697132589

Precision: 0.9716056485978953

Recall: 0.9706566697132589

F1 Score: 0.9709798334183373

AUC: 0.9595745868222374



Metrics by Label (SMOTE):

	Label	Accuracy	Method
0	Benign	0.976393	SMOTE
1	DDoS attack-HOIC	0.963017	SMOTE
2	DDoS attacks-LOIC-HTTP	0.984310	SMOTE
3	DoS attacks-Hulk	0.996869	SMOTE
4	DoS attacks-SlowHTTPTest	1.000000	SMOTE
5	SSH-Bruteforce	0.999467	SMOTE
6	FTP-BruteForce	1.000000	SMOTE
7	Infiltration	0.247962	SMOTE
8	DoS attacks-GoldenEye	0.998814	SMOTE
9	Bot	0.996349	SMOTE
10	DDoS attack-LOIC-UDP	1.000000	SMOTE
11	DoS attacks-Slowloris	0.981043	SMOTE
12	Brute Force -Web	0.636364	SMOTE
13	Brute Force -XSS	1.000000	SMOTE
14	SQL Injection	0.500000	SMOTE

```
[32]: decision_tree_model_adasyn = DecisionTreeClassifier()  
decision_tree_model_adasyn.fit(scaler_adasyn.transform(X_adasyn_train),  
↪Y_adasyn_train.is_attack)
```

```
[32]: DecisionTreeClassifier()
```

```
[33]: # Predict and evaluate on the test set  
dt_metrics["adasyn"] = test_metrics("Decision Tree",  
↪decision_tree_model_adasyn, "ADASYN", scaler_adasyn)
```

Decision Tree with ADASYN Test Set Performance

Classification Report (Test Decision Tree (ADASYN)):

	precision	recall	f1-score	support
0	0.9871	0.9770	0.9820	267720
1	0.8936	0.9378	0.9152	55045
accuracy			0.9704	322765
macro avg	0.9403	0.9574	0.9486	322765
weighted avg	0.9711	0.9704	0.9706	322765

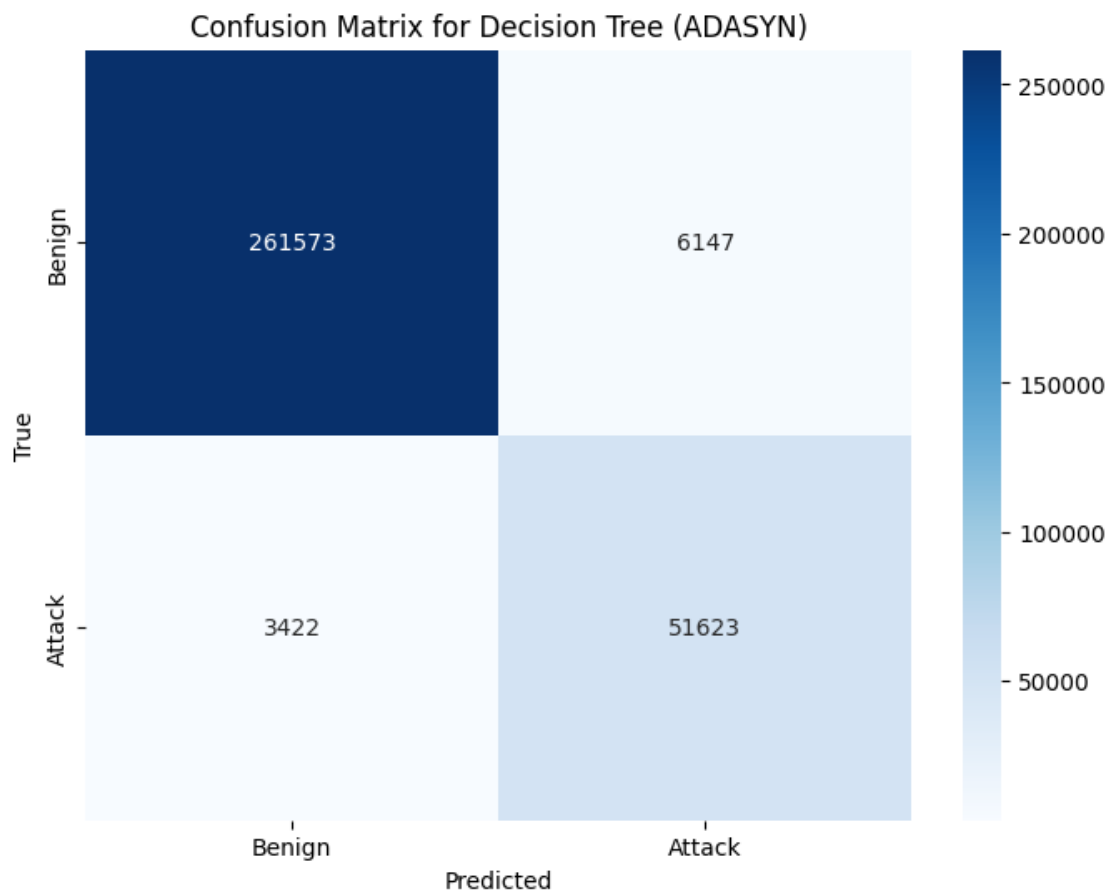
Accuracy: 0.9703530432357907

Precision: 0.9711423595361217

Recall: 0.9703530432357907

F1 Score: 0.9706352716478273

AUC: 0.9574360632731159



Metrics by Label (ADASYN):


	Label	Accuracy	Method
0	Benign	0.977039	ADASYN
1	DDoS attack-HOIC	0.965124	ADASYN
2	DDoS attacks-LOIC-HTTP	0.991505	ADASYN
3	DoS attacks-Hulk	0.997516	ADASYN
4	DoS attacks-SlowHTTPTest	1.000000	ADASYN
5	SSH-Bruteforce	0.999467	ADASYN
6	FTP-BruteForce	1.000000	ADASYN
7	Infiltration	0.124451	ADASYN
8	DoS attacks-GoldenEye	0.997628	ADASYN
9	Bot	0.998088	ADASYN
10	DDoS attack-LOIC-UDP	1.000000	ADASYN
11	DoS attacks-Slowloris	0.971564	ADASYN
12	Brute Force -Web	0.545455	ADASYN
13	Brute Force -XSS	0.800000	ADASYN
14	SQL Injection	0.500000	ADASYN

```
[34]: # Combine metrics into one DataFrame for Decision Tree
combined_metrics_dt = pd.concat([dt_metrics["adasyn"][1],
    ↪ dt_metrics["original"][1], dt_metrics["random"][1], dt_metrics["smote"][1]])
# Pivot the table to get accuracy for each method as columns in the specified
    ↪ order
accuracy_pivot_dt = combined_metrics_dt.pivot(index='Label', columns='Method',
    ↪ values='Accuracy')
accuracy_pivot_dt = accuracy_pivot_dt[['Original', 'Random Oversampling',
    ↪ 'SMOTE', 'ADASYN']]
print("Accuracy by Label and Method (Decision Tree):")
print(accuracy_pivot_dt)
```

Accuracy by Label and Method (Decision Tree):

Method	Original	Random Oversampling	SMOTE	ADASYN
Label				
Benign	0.986602	0.972423	0.976393	0.977039
Bot	0.994958	0.997392	0.996349	0.998088
Brute Force -Web	0.545455	0.545455	0.636364	0.545455
Brute Force -XSS	0.400000	0.600000	1.000000	0.800000
DDoS attack-HOIC	0.951609	0.967812	0.963017	0.965124
DDoS attack-LOIC-UDP	1.000000	1.000000	1.000000	1.000000
DDoS attacks-LOIC-HTTP	0.981016	0.981276	0.984310	0.991505
DoS attacks-GoldenEye	0.998814	1.000000	0.998814	0.997628
DoS attacks-Hulk	0.995681	0.997408	0.996869	0.997516
DoS attacks-SlowHTTPTest	1.000000	1.000000	1.000000	1.000000
DoS attacks-Slowloris	0.962085	0.952607	0.981043	0.971564
FTP-BruteForce	1.000000	1.000000	1.000000	1.000000
Infiltration	0.101881	0.239498	0.247962	0.124451
SQL Injection	0.500000	1.000000	0.500000	0.500000
SSH-Bruteforce	0.999467	0.999733	0.999467	0.999467

```
[35]: plot_overall_accuracy(dt_metrics)
```



1.4.5 Random Forest

```
[36]: rf_metrics = {}
```

```
[37]: rf_model = RandomForestClassifier(verbose=1, n_jobs=-1)
      rf_model.fit(scaler.transform(X_train), Y_train.is_attack)
```

```
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 16 concurrent
workers.
```

```
[Parallel(n_jobs=-1)]: Done 18 tasks      | elapsed: 22.7s
```

```
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed: 1.3min finished
```

```
[37]: RandomForestClassifier(n_jobs=-1, verbose=1)
```

```
[38]: # Predict and evaluate on the test set
      rf_metrics["original"] = test_metrics("Random Forest", rf_model, "Original",
      ↪scaler)
```

```
[Parallel(n_jobs=16)]: Using backend ThreadingBackend with 16 concurrent
workers.
```

Random Forest with Original Test Set Performance

```
[Parallel(n_jobs=16)]: Done 18 tasks      | elapsed: 0.1s
```

```
[Parallel(n_jobs=16)]: Done 100 out of 100 | elapsed: 0.6s finished
```

Classification Report (Test Random Forest (Original)):

	precision	recall	f1-score	support
0	0.9874	0.9920	0.9897	267720
1	0.9600	0.9383	0.9490	55045
accuracy			0.9828	322765
macro avg	0.9737	0.9651	0.9693	322765
weighted avg	0.9827	0.9828	0.9827	322765


Accuracy: 0.9828110235000697

Precision: 0.9827056123070539

Recall: 0.9828110235000697

F1 Score: 0.9827328464227976

AUC: 0.9651406029773095



Metrics by Label (Original):

	Label	Accuracy	Method
0	Benign	0.991958	Original

1	DDoS attack-HOIC	0.989029	Original
2	DDoS attacks-LOIC-HTTP	0.986997	Original
3	DoS attacks-Hulk	0.998056	Original
4	DoS attacks-SlowHTTPTest	1.000000	Original
5	SSH-Bruteforce	0.999467	Original
6	FTP-BruteForce	1.000000	Original
7	Infiltration	0.044828	Original
8	DoS attacks-GoldenEye	1.000000	Original
9	Bot	0.997914	Original
10	DDoS attack-LOIC-UDP	1.000000	Original
11	DoS attacks-Slowloris	0.971564	Original
12	Brute Force -Web	0.545455	Original
13	Brute Force -XSS	0.400000	Original
14	SQL Injection	0.500000	Original

```
[39]: rf_model_random = RandomForestClassifier(verbose=1, n_jobs=-1)
      rf_model_random.fit(scaler_random.transform(X_random_train), Y_random_train.
      ↪is_attack)
```

[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 16 concurrent workers.

[Parallel(n_jobs=-1)]: Done 18 tasks | elapsed: 35.1s

[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed: 2.0min finished

```
[39]: RandomForestClassifier(n_jobs=-1, verbose=1)
```

```
[40]: # Predict and evaluate on the test set
      rf_metrics["random"] = test_metrics("Random Forest", rf_model_random, "Random_
      ↪Oversampling", scaler_random)
```

Random Forest with Random Oversampling Test Set Performance

[Parallel(n_jobs=16)]: Using backend ThreadingBackend with 16 concurrent workers.

[Parallel(n_jobs=16)]: Done 18 tasks | elapsed: 0.1s

[Parallel(n_jobs=16)]: Done 100 out of 100 | elapsed: 0.6s finished

Classification Report (Test Random Forest (Random Oversampling)):

	precision	recall	f1-score	support
0	0.9897	0.9759	0.9828	267720
1	0.8901	0.9508	0.9194	55045
accuracy			0.9716	322765
macro avg	0.9399	0.9633	0.9511	322765
weighted avg	0.9727	0.9716	0.9720	322765


Accuracy: 0.9715861385218347

Precision: 0.972747152441986

Recall: 0.9715861385218347

F1 Score: 0.9719547712798587

AUC: 0.9633459288372468



Metrics by Label (Random Oversampling):

	Label	Accuracy	Method
0	Benign	0.975852	Random Oversampling
1	DDoS attack-HOIC	0.993388	Random Oversampling
2	DDoS attacks-LOIC-HTTP	0.988211	Random Oversampling
3	DoS attacks-Hulk	0.998380	Random Oversampling
4	DoS attacks-SlowHTTPTest	1.000000	Random Oversampling
5	SSH-Bruteforce	0.999467	Random Oversampling
6	FTP-BruteForce	1.000000	Random Oversampling
7	Infiltration	0.232915	Random Oversampling
8	DoS attacks-GoldenEye	1.000000	Random Oversampling
9	Bot	0.999478	Random Oversampling
10	DDoS attack-LOIC-UDP	1.000000	Random Oversampling
11	DoS attacks-Slowloris	0.976303	Random Oversampling
12	Brute Force -Web	0.545455	Random Oversampling
13	Brute Force -XSS	0.600000	Random Oversampling

```
[41]: rf_model_smote = RandomForestClassifier(verbose=1, n_jobs=-1)
      rf_model_smote.fit(scaler_smote.transform(X_smote_train), Y_smote_train.
      ↪is_attack)
```

[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 16 concurrent workers.

[Parallel(n_jobs=-1)]: Done 18 tasks | elapsed: 37.0s

[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed: 2.1min finished

```
[41]: RandomForestClassifier(n_jobs=-1, verbose=1)
```

```
[42]: # Predict and evaluate on the test set
      rf_metrics["smote"] = test_metrics("Random Forest", rf_model_smote, "SMOTE",
      ↪scaler_smote)
```

Random Forest with SMOTE Test Set Performance

[Parallel(n_jobs=16)]: Using backend ThreadingBackend with 16 concurrent workers.

[Parallel(n_jobs=16)]: Done 18 tasks | elapsed: 0.2s

[Parallel(n_jobs=16)]: Done 100 out of 100 | elapsed: 0.8s finished

Classification Report (Test Random Forest (SMOTE)):

	precision	recall	f1-score	support
0	0.9900	0.9737	0.9818	267720
1	0.8817	0.9524	0.9157	55045
accuracy			0.9701	322765
macro avg	0.9359	0.9630	0.9488	322765
weighted avg	0.9716	0.9701	0.9705	322765

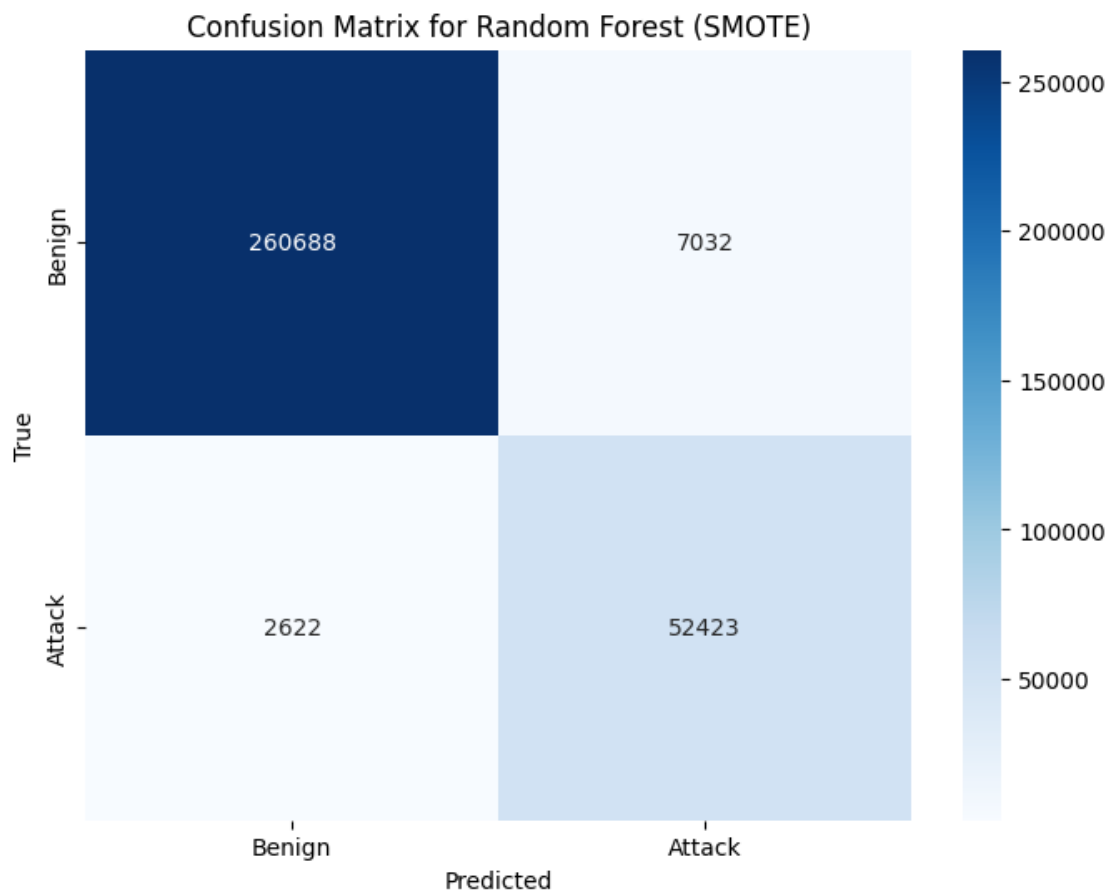
Accuracy: 0.9700896937400276

Precision: 0.9715696425576648

Recall: 0.9700896937400276

F1 Score: 0.9705414989103793

AUC: 0.9630499987398763



Metrics by Label (SMOTE):

	Label	Accuracy	Method
0	Benign	0.973734	SMOTE
1	DDoS attack-HOIC	0.994623	SMOTE
2	DDoS attacks-LOIC-HTTP	0.989511	SMOTE
3	DoS attacks-Hulk	0.998596	SMOTE
4	DoS attacks-SlowHTTPTest	1.000000	SMOTE
5	SSH-Bruteforce	0.999467	SMOTE
6	FTP-BruteForce	1.000000	SMOTE
7	Infiltration	0.247962	SMOTE
8	DoS attacks-GoldenEye	1.000000	SMOTE
9	Bot	0.999478	SMOTE
10	DDoS attack-LOIC-UDP	1.000000	SMOTE
11	DoS attacks-Slowloris	0.981043	SMOTE
12	Brute Force -Web	0.727273	SMOTE
13	Brute Force -XSS	0.600000	SMOTE
14	SQL Injection	0.500000	SMOTE

```
[43]: rf_model_adasyn = RandomForestClassifier(verbose=1, n_jobs=-1)
      rf_model_adasyn.fit(scaler_adasyn.transform(X_adasyn_train), Y_adasyn_train.
      ↪is_attack)
```

[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 16 concurrent workers.

[Parallel(n_jobs=-1)]: Done 18 tasks | elapsed: 37.5s

[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed: 2.1min finished

```
[43]: RandomForestClassifier(n_jobs=-1, verbose=1)
```

```
[44]: # Predict and evaluate on the test set
      rf_metrics["adasyn"] = test_metrics("Random Forest", rf_model_adasyn, "ADASYN",
      ↪scaler_adasyn)
```

Random Forest with ADASYN Test Set Performance

[Parallel(n_jobs=16)]: Using backend ThreadingBackend with 16 concurrent workers.

[Parallel(n_jobs=16)]: Done 18 tasks | elapsed: 0.1s

[Parallel(n_jobs=16)]: Done 100 out of 100 | elapsed: 0.7s finished

Classification Report (Test Random Forest (ADASYN)):

	precision	recall	f1-score	support
0	0.9889	0.9315	0.9593	267720
1	0.7402	0.9490	0.8317	55045
accuracy			0.9345	322765
macro avg	0.8646	0.9403	0.8955	322765
weighted avg	0.9465	0.9345	0.9376	322765


Accuracy: 0.934515824206466

Precision: 0.9464757098207695

Recall: 0.934515824206466

F1 Score: 0.9375849455115585

AUC: 0.9402854513910675



Metrics by Label (ADASYN):


	Label	Accuracy	Method
0	Benign	0.931529	ADASYN
1	DDoS attack-HOIC	0.997748	ADASYN
2	DDoS attacks-LOIC-HTTP	0.995232	ADASYN
3	DoS attacks-Hulk	0.999568	ADASYN
4	DoS attacks-SlowHTTPTest	1.000000	ADASYN
5	SSH-Bruteforce	0.999733	ADASYN
6	FTP-BruteForce	1.000000	ADASYN
7	Infiltration	0.151411	ADASYN
8	DoS attacks-GoldenEye	1.000000	ADASYN
9	Bot	0.999478	ADASYN
10	DDoS attack-LOIC-UDP	1.000000	ADASYN
11	DoS attacks-Slowloris	0.990521	ADASYN
12	Brute Force -Web	0.909091	ADASYN
13	Brute Force -XSS	0.800000	ADASYN
14	SQL Injection	1.000000	ADASYN

```
[45]: # Combine metrics into one DataFrame
combined_metrics_rf = pd.concat([rf_metrics["adasyn"][1],
    ↪ rf_metrics["original"][1], rf_metrics["random"][1], rf_metrics["smote"][1]])
# Pivot the table to get accuracy for each method as columns in the specified
    ↪ order
accuracy_pivot_rf = combined_metrics_rf.pivot(index='Label', columns='Method',
    ↪ values='Accuracy')
accuracy_pivot_rf = accuracy_pivot_rf[['Original', 'Random Oversampling',
    ↪ 'SMOTE', 'ADASYN']]
print("Accuracy by Label and Method:")
print(accuracy_pivot_rf)
```

Accuracy by Label and Method:

Method	Original	Random Oversampling	SMOTE	ADASYN
Label				
Benign	0.991958	0.975852	0.973734	0.931529
Bot	0.997914	0.999478	0.999478	0.999478
Brute Force -Web	0.545455	0.545455	0.727273	0.909091
Brute Force -XSS	0.400000	0.600000	0.600000	0.800000
DDoS attack-HOIC	0.989029	0.993388	0.994623	0.997748
DDoS attack-LOIC-UDP	1.000000	1.000000	1.000000	1.000000
DDoS attacks-LOIC-HTTP	0.986997	0.988211	0.989511	0.995232
DoS attacks-GoldenEye	1.000000	1.000000	1.000000	1.000000
DoS attacks-Hulk	0.998056	0.998380	0.998596	0.999568
DoS attacks-SlowHTTPTest	1.000000	1.000000	1.000000	1.000000
DoS attacks-Slowloris	0.971564	0.976303	0.981043	0.990521
FTP-BruteForce	1.000000	1.000000	1.000000	1.000000
Infiltration	0.044828	0.232915	0.247962	0.151411
SQL Injection	0.500000	1.000000	0.500000	1.000000
SSH-Bruteforce	0.999467	0.999467	0.999467	0.999733

```
[46]: plot_overall_accuracy(rf_metrics)
```



1.4.6 Adaboost

```
[47]: ada_metrics = {}
```

```
[48]: from sklearn.ensemble import AdaBoostClassifier

ada_model = AdaBoostClassifier(algorithm='SAMME')
ada_model.fit(scaler.transform(X_train), Y_train.is_attack)
```

```
[48]: AdaBoostClassifier(algorithm='SAMME')
```

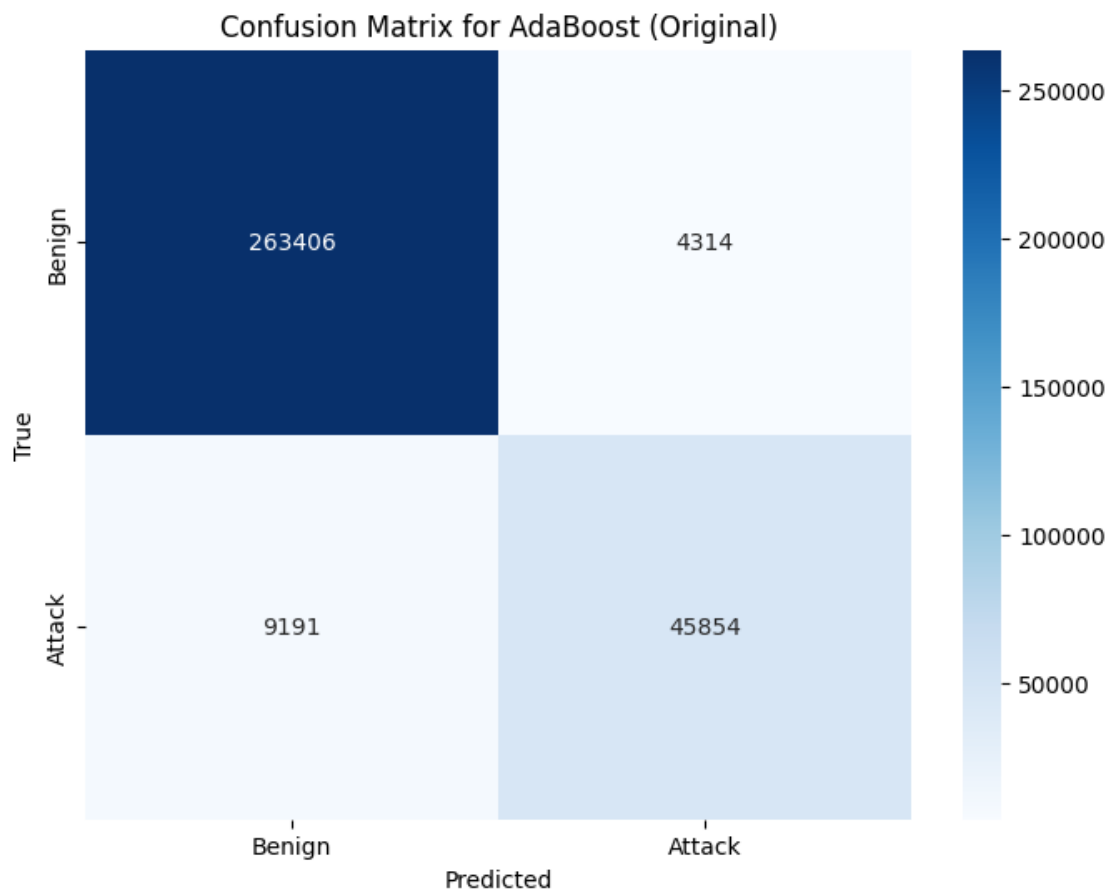
```
[49]: # Predict and evaluate on the test set
ada_metrics["original"] = test_metrics("AdaBoost", ada_model, "Original",
↪scaler)
```

AdaBoost with Original Test Set Performance

Classification Report (Test AdaBoost (Original)):

	precision	recall	f1-score	support
0	0.9663	0.9839	0.9750	267720
1	0.9140	0.8330	0.8716	55045
accuracy			0.9582	322765
macro avg	0.9401	0.9085	0.9233	322765
weighted avg	0.9574	0.9582	0.9574	322765

Accuracy: 0.958158412467275
Precision: 0.9573685394091846
Recall: 0.958158412467275
F1 Score: 0.9573774935315756
AUC: 0.9084568363222154



Metrics by Label (Original):

	Label	Accuracy	Method
0	Benign	0.983886	Original
1	DDoS attack-HOIC	0.957640	Original
2	DDoS attacks-LOIC-HTTP	0.923544	Original
3	DoS attacks-Hulk	0.974949	Original
4	DoS attacks-SlowHTTPTest	0.869922	Original
5	SSH-Bruteforce	0.980011	Original
6	FTP-BruteForce	0.874452	Original
7	Infiltration	0.036991	Original
8	DoS attacks-GoldenEye	0.501779	Original
9	Bot	0.492524	Original

10	DDoS attack-LOIC-UDP	0.000000	Original
11	DoS attacks-Slowloris	0.497630	Original
12	Brute Force -Web	0.090909	Original
13	Brute Force -XSS	0.200000	Original
14	SQL Injection	0.000000	Original

```
[50]: ada_model_random = AdaBoostClassifier(algorithm='SAMME')
      ada_model_random.fit(scaler_random.transform(X_random_train), Y_random_train.
      ↪is_attack)
```

```
[50]: AdaBoostClassifier(algorithm='SAMME')
```

```
[51]: # Predict and evaluate on the test set
      ada_metrics["random"] = test_metrics("AdaBoost", ada_model_random, "Random_
      ↪Oversampling", scaler_random)
```

AdaBoost with Random Oversampling Test Set Performance

Classification Report (Test AdaBoost (Random Oversampling)):

	precision	recall	f1-score	support
0	0.9897	0.8366	0.9068	267720
1	0.5466	0.9577	0.6960	55045
accuracy			0.8573	322765
macro avg	0.7681	0.8972	0.8014	322765
weighted avg	0.9141	0.8573	0.8708	322765


Accuracy: 0.8572831626725327

Precision: 0.9141432761392969

Recall: 0.8572831626725327

F1 Score: 0.8708068473290808

AUC: 0.8971857421926237



Metrics by Label (Random Oversampling):

	Label	Accuracy	Method
0	Benign	0.836628	Random Oversampling
1	DDoS attack-HOIC	1.000000	Random Oversampling
2	DDoS attacks-LOIC-HTTP	0.986304	Random Oversampling
3	DoS attacks-Hulk	0.999136	Random Oversampling
4	DoS attacks-SlowHTTPTest	1.000000	Random Oversampling
5	SSH-Bruteforce	1.000000	Random Oversampling
6	FTP-BruteForce	1.000000	Random Oversampling
7	Infiltration	0.331034	Random Oversampling
8	DoS attacks-GoldenEye	1.000000	Random Oversampling
9	Bot	0.995480	Random Oversampling
10	DDoS attack-LOIC-UDP	1.000000	Random Oversampling
11	DoS attacks-Slowloris	1.000000	Random Oversampling
12	Brute Force -Web	1.000000	Random Oversampling
13	Brute Force -XSS	1.000000	Random Oversampling
14	SQL Injection	1.000000	Random Oversampling

```
[52]: ada_model_smote = AdaBoostClassifier(algorithm='SAMME')
ada_model_smote.fit(scaler_smote.transform(X_smote_train), Y_smote_train.
↳is_attack)
```

```
[52]: AdaBoostClassifier(algorithm='SAMME')
```

```
[53]: # Predict and evaluate on the test set
ada_metrics["smote"] = test_metrics("AdaBoost", ada_model_smote, "SMOTE",
↳scaler_smote)
```

AdaBoost with SMOTE Test Set Performance

Classification Report (Test AdaBoost (SMOTE)):

	precision	recall	f1-score	support
0	0.9892	0.8172	0.8950	267720
1	0.5183	0.9567	0.6723	55045
accuracy			0.8410	322765
macro avg	0.7537	0.8869	0.7837	322765
weighted avg	0.9089	0.8410	0.8570	322765

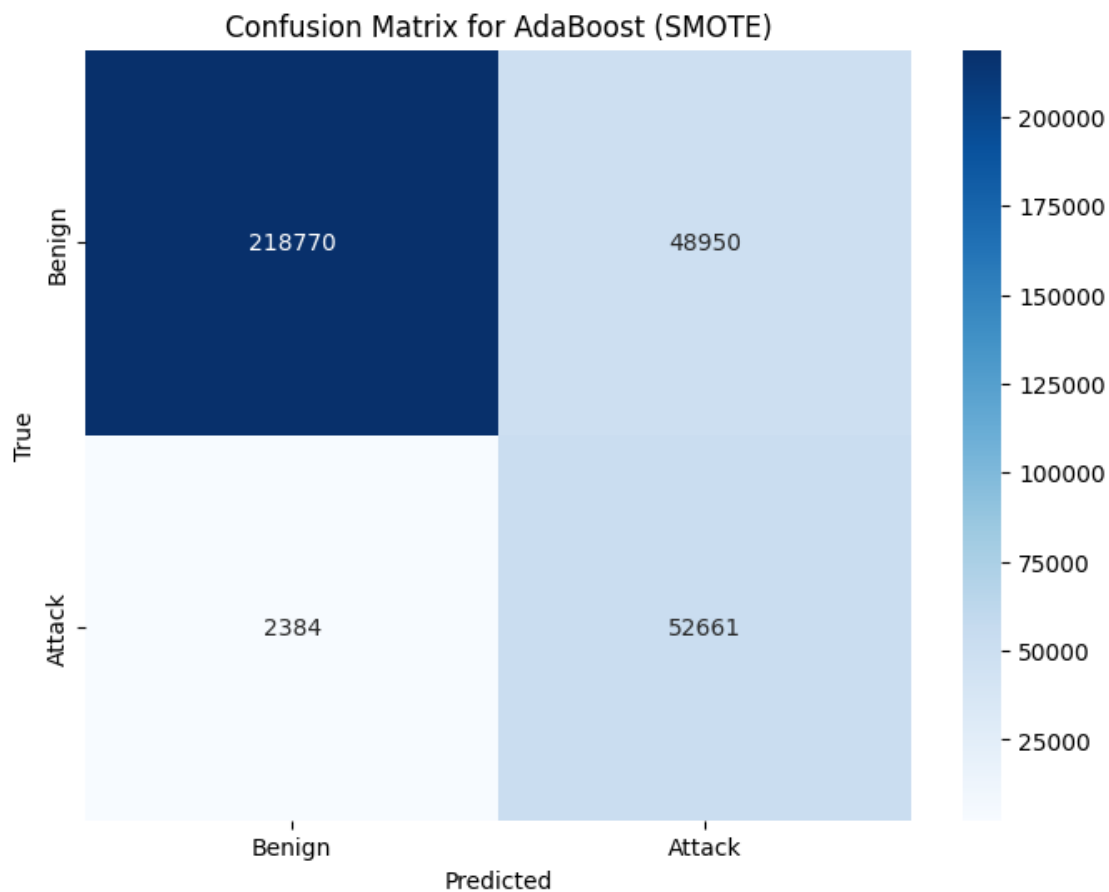
Accuracy: 0.8409554939352161

Precision: 0.9089018131038703

Recall: 0.8409554939352161

F1 Score: 0.8570188688361758

AUC: 0.8869248500171077



Metrics by Label (SMOTE):

	Label	Accuracy	Method
0	Benign	0.817160	SMOTE
1	DDoS attack-HOIC	1.000000	SMOTE
2	DDoS attacks-LOIC-HTTP	0.986564	SMOTE
3	DoS attacks-Hulk	0.999784	SMOTE
4	DoS attacks-SlowHTTPTest	1.000000	SMOTE
5	SSH-Bruteforce	1.000000	SMOTE
6	FTP-BruteForce	1.000000	SMOTE
7	Infiltration	0.334483	SMOTE
8	DoS attacks-GoldenEye	1.000000	SMOTE
9	Bot	0.981919	SMOTE
10	DDoS attack-LOIC-UDP	1.000000	SMOTE
11	DoS attacks-Slowloris	1.000000	SMOTE
12	Brute Force -Web	1.000000	SMOTE
13	Brute Force -XSS	1.000000	SMOTE
14	SQL Injection	1.000000	SMOTE

```
[54]: ada_model_adasyn = AdaBoostClassifier(algorithm='SAMME')
ada_model_adasyn.fit(scaler_adasyn.transform(X_adasyn_train), Y_adasyn_train.
↳is_attack)
```

```
[54]: AdaBoostClassifier(algorithm='SAMME')
```

```
[55]: # Predict and evaluate on the test set
ada_metrics["adasyn"] = test_metrics("AdaBoost", ada_model_adasyn, "ADASYN",
↳scaler_adasyn)
```

AdaBoost with ADASYN Test Set Performance

Classification Report (Test AdaBoost (ADASYN)):

	precision	recall	f1-score	support
0	0.9529	0.8393	0.8925	267720
1	0.5053	0.7981	0.6188	55045
accuracy			0.8323	322765
macro avg	0.7291	0.8187	0.7557	322765
weighted avg	0.8765	0.8323	0.8458	322765

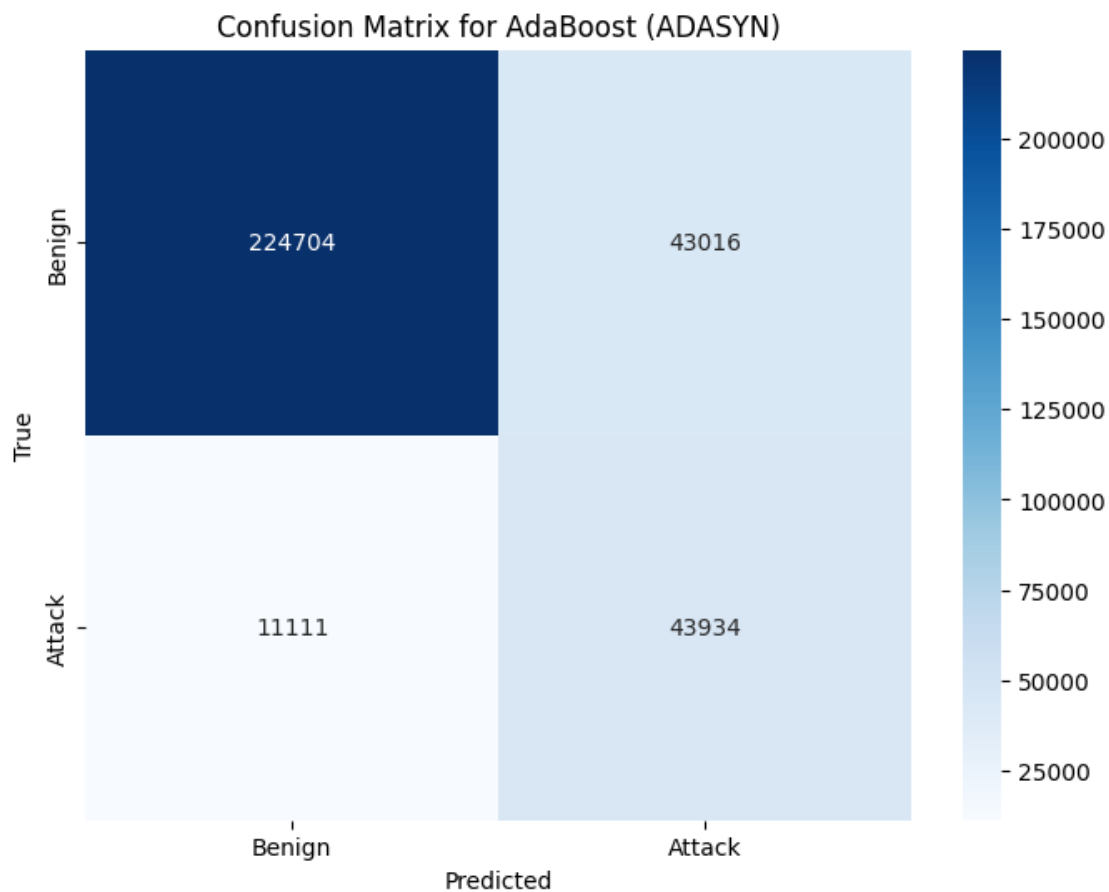
Accuracy: 0.8323021393273744

Precision: 0.876547317034483

Recall: 0.8323021393273744

F1 Score: 0.8458294017160278

AUC: 0.8187358191117472



Metrics by Label (ADASYN):

	Label	Accuracy	Method
0	Benign	0.839325	ADASYN
1	DDoS attack-HOIC	0.999419	ADASYN
2	DDoS attacks-LOIC-HTTP	0.481276	ADASYN
3	DoS attacks-Hulk	0.996977	ADASYN
4	DoS attacks-SlowHTTPTest	1.000000	ADASYN
5	SSH-Bruteforce	1.000000	ADASYN
6	FTP-BruteForce	1.000000	ADASYN
7	Infiltration	0.318495	ADASYN
8	DoS attacks-GoldenEye	1.000000	ADASYN
9	Bot	0.493220	ADASYN
10	DDoS attack-LOIC-UDP	1.000000	ADASYN
11	DoS attacks-Slowloris	1.000000	ADASYN
12	Brute Force -Web	0.818182	ADASYN
13	Brute Force -XSS	1.000000	ADASYN
14	SQL Injection	1.000000	ADASYN

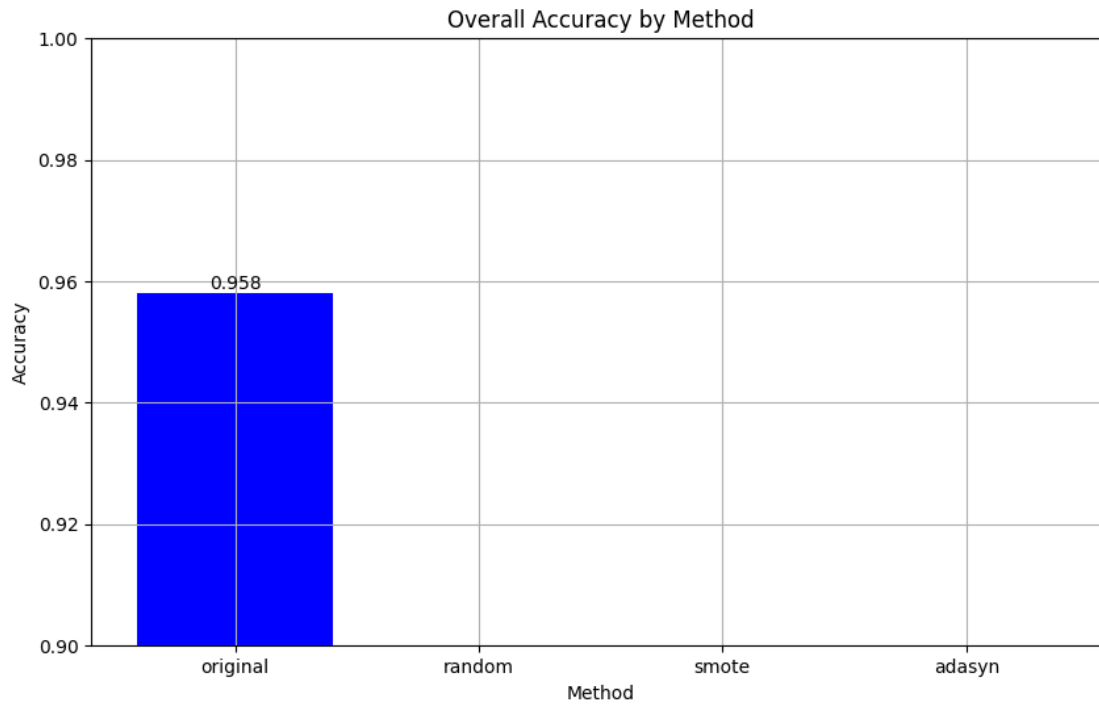
```
[56]: # Combine metrics into one DataFrame
combined_metrics_ada = pd.concat([ada_metrics["adasyn"][1],
    ↪ada_metrics["original"][1], ada_metrics["random"][1],
    ↪ada_metrics["smote"][1]])

# Pivot the table to get accuracy for each method as columns in the specified
    ↪order
accuracy_pivot_ada = combined_metrics_ada.pivot(index='Label',
    ↪columns='Method', values='Accuracy')
accuracy_pivot_ada = accuracy_pivot_ada[['Original', 'Random Oversampling',
    ↪'SMOTE', 'ADASYN']]
print("Accuracy by Label and Method (AdaBoost):")
print(accuracy_pivot_ada)
```

Accuracy by Label and Method (AdaBoost):

Method	Original	Random Oversampling	SMOTE	ADASYN
Label				
Benign	0.983886	0.836628	0.817160	0.839325
Bot	0.492524	0.995480	0.981919	0.493220
Brute Force -Web	0.090909	1.000000	1.000000	0.818182
Brute Force -XSS	0.200000	1.000000	1.000000	1.000000
DDOS attack-HOIC	0.957640	1.000000	1.000000	0.999419
DDOS attack-LOIC-UDP	0.000000	1.000000	1.000000	1.000000
DDoS attacks-LOIC-HTTP	0.923544	0.986304	0.986564	0.481276
DoS attacks-GoldenEye	0.501779	1.000000	1.000000	1.000000
DoS attacks-Hulk	0.974949	0.999136	0.999784	0.996977
DoS attacks-SlowHTTPTest	0.869922	1.000000	1.000000	1.000000
DoS attacks-Slowloris	0.497630	1.000000	1.000000	1.000000
FTP-BruteForce	0.874452	1.000000	1.000000	1.000000
Infiltration	0.036991	0.331034	0.334483	0.318495
SQL Injection	0.000000	1.000000	1.000000	1.000000
SSH-Bruteforce	0.980011	1.000000	1.000000	1.000000

```
[57]: plot_overall_accuracy(ada_metrics)
```



1.4.7 XGBoost

```
[58]: xgb_metrics = {}
```

```
[59]: import xgboost as xgb
```

```
xgb_model = xgb.XGBClassifier(n_jobs=-1)
xgb_model.fit(scaler.transform(X_train), Y_train.is_attack)
```

```
[59]: XGBClassifier(base_score=None, booster=None, callbacks=None,
                  colsample_bylevel=None, colsample_bynode=None,
                  colsample_bytree=None, device=None, early_stopping_rounds=None,
                  enable_categorical=False, eval_metric=None, feature_types=None,
                  gamma=None, grow_policy=None, importance_type=None,
```



```

interaction_constraints=None, learning_rate=None, max_bin=None,
max_cat_threshold=None, max_cat_to_onehot=None,
max_delta_step=None, max_depth=None, max_leaves=None,
min_child_weight=None, missing=nan, monotone_constraints=None,
multi_strategy=None, n_estimators=None, n_jobs=-1,
num_parallel_tree=None, random_state=None, ...)

```

```

[60]: # Predict and evaluate on the test set
      # Original Dataset
      xgb_metrics["original"] = test_metrics("XGBoost", xgb_model, "Original", scaler)

```

XGBoost with Original Test Set Performance

Classification Report (Test XGBoost (Original)):

	precision	recall	f1-score	support
0	0.9872	0.9933	0.9902	267720
1	0.9665	0.9372	0.9517	55045
accuracy			0.9838	322765
macro avg	0.9769	0.9653	0.9710	322765
weighted avg	0.9837	0.9838	0.9837	322765

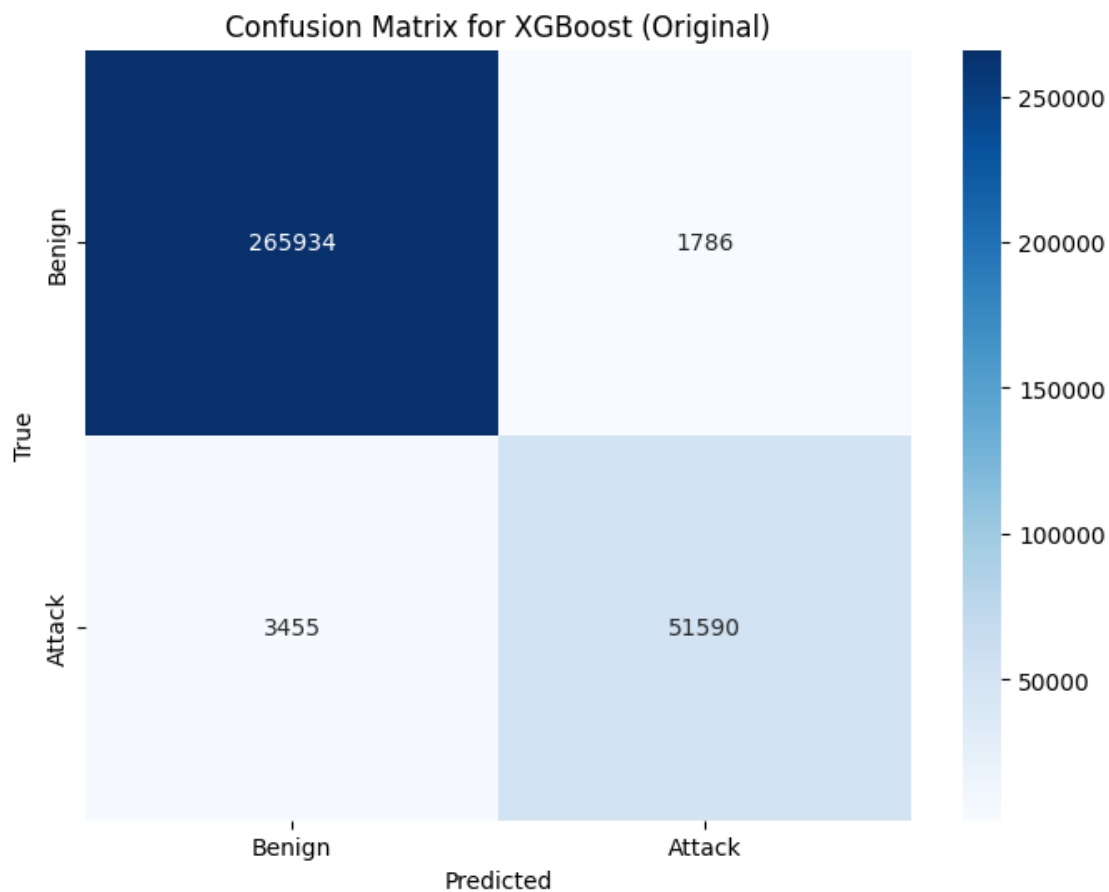
Accuracy: 0.983762179914179

Precision: 0.983655475338238

Recall: 0.983762179914179

F1 Score: 0.9836624283867066

AUC: 0.9652810119484841



Metrics by Label (Original):

	Label	Accuracy	Method
0	Benign	0.993329	Original
1	DDoS attack-HOIC	0.996440	Original
2	DDoS attacks-LOIC-HTTP	0.979282	Original
3	DoS attacks-Hulk	0.999676	Original
4	DoS attacks-SlowHTTPTest	1.000000	Original
5	SSH-Bruteforce	1.000000	Original
6	FTP-BruteForce	1.000000	Original
7	Infiltration	0.023511	Original
8	DoS attacks-GoldenEye	0.998814	Original
9	Bot	0.996871	Original
10	DDoS attack-LOIC-UDP	1.000000	Original
11	DoS attacks-Slowloris	0.914692	Original
12	Brute Force -Web	0.363636	Original
13	Brute Force -XSS	0.400000	Original
14	SQL Injection	0.000000	Original

```
[61]: xgb_model_random = xgb.XGBClassifier(n_jobs=-1)
xgb_model_random.fit(scaler_random.transform(X_random_train), Y_random_train,
↳ is_attack)
```

```
[61]: XGBClassifier(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=None, device=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric=None, feature_types=None,
    gamma=None, grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=None, max_bin=None,
    max_cat_threshold=None, max_cat_to_onehot=None,
    max_delta_step=None, max_depth=None, max_leaves=None,
    min_child_weight=None, missing=nan, monotone_constraints=None,
    multi_strategy=None, n_estimators=None, n_jobs=-1,
    num_parallel_tree=None, random_state=None, ...)
```

```
[62]: # Predict and evaluate on the test set
# Random Oversampling
xgb_metrics["random"] = test_metrics("XGBoost", xgb_model_random, "Random_↳
↳ Oversampling", scaler_random)
```

XGBoost with Random Oversampling Test Set Performance

Classification Report (Test XGBoost (Random Oversampling)):

	precision	recall	f1-score	support
0	0.9902	0.9752	0.9826	267720
1	0.8876	0.9531	0.9192	55045
accuracy			0.9714	322765
macro avg	0.9389	0.9641	0.9509	322765
weighted avg	0.9727	0.9714	0.9718	322765


Accuracy: 0.9714126376775672

Precision: 0.9727028383818523

Recall: 0.9714126376775672

F1 Score: 0.9718117201171353

AUC: 0.9641216761758172



Metrics by Label (Random Oversampling):

	Label	Accuracy	Method
0	Benign	0.975187	Random Oversampling
1	DDoS attack-HOIC	0.998329	Random Oversampling
2	DDoS attacks-LOIC-HTTP	0.985697	Random Oversampling
3	DoS attacks-Hulk	0.999784	Random Oversampling
4	DoS attacks-SlowHTTPTest	1.000000	Random Oversampling
5	SSH-Bruteforce	1.000000	Random Oversampling
6	FTP-BruteForce	1.000000	Random Oversampling
7	Infiltration	0.254859	Random Oversampling
8	DoS attacks-GoldenEye	1.000000	Random Oversampling
9	Bot	0.998435	Random Oversampling
10	DDoS attack-LOIC-UDP	1.000000	Random Oversampling
11	DoS attacks-Slowloris	0.985782	Random Oversampling
12	Brute Force -Web	0.636364	Random Oversampling
13	Brute Force -XSS	0.800000	Random Oversampling
14	SQL Injection	1.000000	Random Oversampling

```
[63]: xgb_model_smote = xgb.XGBClassifier(n_jobs=-1)
xgb_model_smote.fit(scaler_smote.transform(X_smote_train), Y_smote_train.
↳is_attack)
```

```
[63]: XGBClassifier(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=None, device=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric=None, feature_types=None,
    gamma=None, grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=None, max_bin=None,
    max_cat_threshold=None, max_cat_to_onehot=None,
    max_delta_step=None, max_depth=None, max_leaves=None,
    min_child_weight=None, missing=nan, monotone_constraints=None,
    multi_strategy=None, n_estimators=None, n_jobs=-1,
    num_parallel_tree=None, random_state=None, ...)
```

```
[64]: # Predict and evaluate on the test set
# SMOTE
xgb_metrics["smote"] = test_metrics("XGBoost", xgb_model_smote, "SMOTE",
↳scaler_smote)
```

XGBoost with SMOTE Test Set Performance

Classification Report (Test XGBoost (SMOTE)):

	precision	recall	f1-score	support
0	0.9899	0.9828	0.9863	267720
1	0.9193	0.9512	0.9350	55045
accuracy			0.9774	322765
macro avg	0.9546	0.9670	0.9607	322765
weighted avg	0.9779	0.9774	0.9776	322765


Accuracy: 0.9774324973277771

Precision: 0.9778513988563949

Recall: 0.9774324973277771

F1 Score: 0.9775848417524753

AUC: 0.9670216661694708



Metrics by Label (SMOTE):

	Label	Accuracy	Method
0	Benign	0.982822	SMOTE
1	DDoS attack-HOIC	0.999055	SMOTE
2	DDoS attacks-LOIC-HTTP	0.982576	SMOTE
3	DoS attacks-Hulk	0.999784	SMOTE
4	DoS attacks-SlowHTTPTest	1.000000	SMOTE
5	SSH-Bruteforce	1.000000	SMOTE
6	FTP-BruteForce	1.000000	SMOTE
7	Infiltration	0.230721	SMOTE
8	DoS attacks-GoldenEye	1.000000	SMOTE
9	Bot	0.998609	SMOTE
10	DDoS attack-LOIC-UDP	1.000000	SMOTE
11	DoS attacks-Slowloris	0.985782	SMOTE
12	Brute Force -Web	0.727273	SMOTE
13	Brute Force -XSS	1.000000	SMOTE
14	SQL Injection	0.500000	SMOTE

```
[65]: xgb_model_adasyn = xgb.XGBClassifier(n_jobs=-1)
xgb_model_adasyn.fit(scaler_adasyn.transform(X_adasyn_train), Y_adasyn_train,
↳ is_attack)
```

```
[65]: XGBClassifier(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=None, device=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric=None, feature_types=None,
    gamma=None, grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=None, max_bin=None,
    max_cat_threshold=None, max_cat_to_onehot=None,
    max_delta_step=None, max_depth=None, max_leaves=None,
    min_child_weight=None, missing=nan, monotone_constraints=None,
    multi_strategy=None, n_estimators=None, n_jobs=-1,
    num_parallel_tree=None, random_state=None, ...)
```

```
[66]: # Predict and evaluate on the test set
# ADASYN
xgb_metrics["adasyn"] = test_metrics("XGBoost", xgb_model_adasyn, "ADASYN",
↳ scaler_adasyn)
```

XGBoost with ADASYN Test Set Performance

Classification Report (Test XGBoost (ADASYN)):

	precision	recall	f1-score	support
0	0.9892	0.9689	0.9789	267720
1	0.8626	0.9484	0.9034	55045
accuracy			0.9654	322765
macro avg	0.9259	0.9587	0.9412	322765
weighted avg	0.9676	0.9654	0.9661	322765

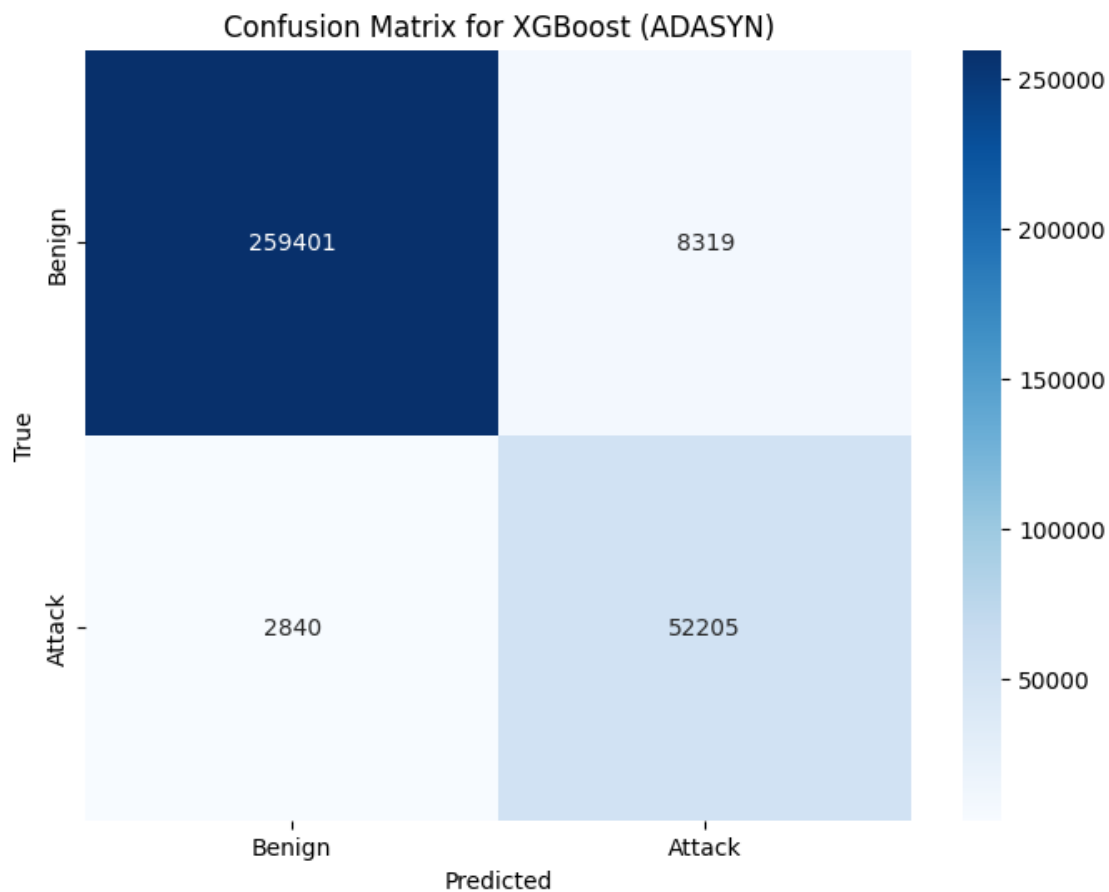
Accuracy: 0.9654268585503385

Precision: 0.9675762560643688

Recall: 0.9654268585503385

F1 Score: 0.966067678825433

AUC: 0.9586661700611769



Metrics by Label (ADASYN):

	Label	Accuracy	Method
0	Benign	0.968926	ADASYN
1	DDoS attack-HOIC	0.999491	ADASYN
2	DDoS attacks-LOIC-HTTP	0.997053	ADASYN
3	DoS attacks-Hulk	1.000000	ADASYN
4	DoS attacks-SlowHTTPTest	1.000000	ADASYN
5	SSH-Bruteforce	1.000000	ADASYN
6	FTP-BruteForce	1.000000	ADASYN
7	Infiltration	0.125392	ADASYN
8	DoS attacks-GoldenEye	0.998814	ADASYN
9	Bot	0.999652	ADASYN
10	DDoS attack-LOIC-UDP	1.000000	ADASYN
11	DoS attacks-Slowloris	0.990521	ADASYN
12	Brute Force -Web	0.727273	ADASYN
13	Brute Force -XSS	1.000000	ADASYN
14	SQL Injection	0.500000	ADASYN


```
[67]: # Combine metrics into one DataFrame
combined_metrics_xgb = pd.concat([xgb_metrics["adasyn"][1],
    ↳xgb_metrics["original"][1], xgb_metrics["random"][1],
    ↳xgb_metrics["smote"][1]])
# Pivot the table to get accuracy for each method as columns in the specified
    ↳order
accuracy_pivot_xgb = combined_metrics_xgb.pivot(index='Label',
    ↳columns='Method', values='Accuracy')
accuracy_pivot_xgb = accuracy_pivot_xgb[['Original', 'Random Oversampling',
    ↳'SMOTE', 'ADASYN']]
print("Accuracy by Label and Method:")
print(accuracy_pivot_xgb)
```

Accuracy by Label and Method:

Method	Original	Random Oversampling	SMOTE	ADASYN
Label				
Benign	0.993329	0.975187	0.982822	0.968926
Bot	0.996871	0.998435	0.998609	0.999652
Brute Force -Web	0.363636	0.636364	0.727273	0.727273
Brute Force -XSS	0.400000	0.800000	1.000000	1.000000
DDoS attack-HOIC	0.996440	0.998329	0.999055	0.999491
DDoS attack-LOIC-UDP	1.000000	1.000000	1.000000	1.000000
DDoS attacks-LOIC-HTTP	0.979282	0.985697	0.982576	0.997053
DoS attacks-GoldenEye	0.998814	1.000000	1.000000	0.998814
DoS attacks-Hulk	0.999676	0.999784	0.999784	1.000000
DoS attacks-SlowHTTPTest	1.000000	1.000000	1.000000	1.000000
DoS attacks-Slowloris	0.914692	0.985782	0.985782	0.990521
FTP-BruteForce	1.000000	1.000000	1.000000	1.000000
Infiltration	0.023511	0.254859	0.230721	0.125392
SQL Injection	0.000000	1.000000	0.500000	0.500000
SSH-Bruteforce	1.000000	1.000000	1.000000	1.000000

```
[68]: plot_overall_accuracy(xgb_metrics)
```

